



# **DATABRICKS STREAM CASE STUDY**




# CASE STUDY



**accenture**<sup>></sup>technology

# NEW YORK CITY TAXI DATA



**accenture**<sup>></sup>**technology**

## Index

1.1	Index .....	4
<b>2</b>	<b>Background .....</b>	<b>5</b>
<b>3</b>	<b>Challenges.....</b>	<b>6</b>
<b>4</b>	<b>Business Need.....</b>	<b>7</b>
<b>5</b>	<b>Proposed Solution .....</b>	<b>8</b>
<b>6</b>	<b>Attributes and Sample Data .....</b>	<b>9</b>

## 2 Background

- New York City has two types of taxis: yellow and green; they are widely recognizable symbols of the city.
- Taxis painted **yellow** (medallion taxis) can pick up passengers from anywhere in the five boroughs.
- Those painted apple **green** (street hail livery vehicles, commonly known as "boro taxis"), which began to appear in August 2013, are allowed to pick up passengers in Upper Manhattan, the Bronx, Brooklyn, Queens (excluding LaGuardia Airport and John F. Kennedy International Airport), and Staten Island.
- Both taxi types have the same fare structure.
- Taxicabs are operated by private companies and licensed by the **New York City Taxi and Limousine Commission (TLC)**.
- It also oversees over 40,000 other **for-hire vehicles (FHVs)**, including "black cars" like Uber, commuter vans, and ambulettes.
- All types of taxis are licensed by the **TLC** which oversees for-hire vehicles, taxis, commuter vans, and paratransit vehicles.
- **Accenture** is responsible for developing and maintaining the data and analytical systems for New York City taxi.

### 3 Challenges

- Things were smooth until the recent arrival of FHV in the scene. Though the system was working well, the challenge started when Accenture started collecting, storing and processing the data for FHV.
- FHV are of multiple types:
  - Community cars, Black cars, Luxury limousines
  - High volume for-hire services which include app-based companies like Uber and Lyft. These dispatches are more than 10,000 trips per day
- The request for FHV by passengers is accepted by bases, and then the bases dispatch the request to the cab drivers.
- There are more than 750 bases and 100000 FHV, and all these different types of FHV operate in different ways.
- The number of FHV are much higher than yellow and green taxis, which led to exponential increase in data volumes.
- The schema for FHV data is a mix of various data formats like CSV, TSV and JSON formats. The sources of this data are quite disparate.

## 4 Business Need

- The requirement is to stage, process and store the data of all types of taxis irrespective of their sources and formats and transform it to the analytical needs.
- Finally build reports/visualizations which provide actionable insights.
- Most of this is needed in real time. There is an increase in demand for stream processing data as well due to higher trip rate.
- The requirement is to ingest and process the data at very high frequencies. Finally, none of the data should be discarded. In fact, it should be preserved for enabling use-cases related to regulatory compliance, passenger safety, insurance, targeted ads/promotions/offers etc.,
- TLC wants Accenture to build a common platform to store all data related to trips, cabs and passengers in order to analyze the data for better business insights like Revenue by taxi type, location, Total trips, max trips by regions etc.,

## 5 Proposed Solution

- Considering the Volume, Velocity and Veracity aspects of the data, Accenture has decided to build a Data Lake using **Databricks Lake House** Platform which is going to be a single store for raw intermittent and processed data.
- Accenture has analyzed various available options and delivered some PoC's. In a one-year contract with TLC, Accenture will build a Spark based Data Lake augmented by a Cloud based Data Lake on Azure.
- The solution will implement the latest features and concepts of **Databricks Lake House** Platform



## 6 Attributes and Sample Data



**Green Taxi**



**Yellow Taxi**

### **Attributes:**

- Pickup & drop time
- Pickup & drop location details
- Trip distance
- Passenger count
- Solo/ shared ride
- Payment info - fare, tips, tolls, tax etc...
- Payment mode



**Green Taxi**

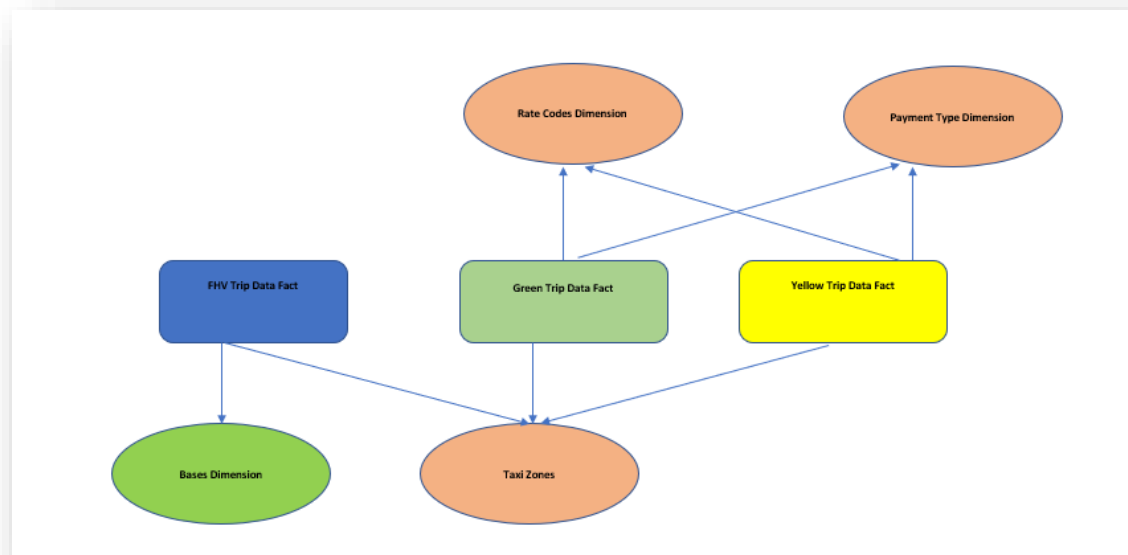


**Yellow Taxi**

### **Attributes:**

- Pickup & drop time
- Pickup & drop location details
- Trip distance
- Passenger count
- Solo/ shared ride
- Payment info - fare, tips, tolls, tax etc...
- Payment mode

- Capture data once a month (in CSV file)
- Extract on-premise using ETL tool
- Build dimensions/facts
- Store it in RDBMS (Data Warehouse)
- Build Aggregated reports & KPI
  - Revenue by taxi type, location etc.,
  - Total trips, max trips by regions etc.,



### Sample Data: fhv\_tripdata\_2019-05.csv/ fhv\_tripdata\_2019-06.csv

dispatching_base_num	B00013
pickup_datetime	2019-06-01 00:51:33
dropoff_datetime	2019-06-01 01:20:07
PULocationID	83
DOLocationID	173
SR_Flag	Null

### Sample Data: FhvBases.json

```

{
  "License Number": "B02865"
  , "Entity Name": "VIER-NY, LLC"
  , "Telephone Number": 6466657536
  , "SHL Endorsed": "No"
  , "Address" :
    {
      "Building": "636"
      , "Street": "WEST 28 STREET"
      , "City": "NEW YORK"
      , "State": "NY"
      , "Postcode": 10001
    }
  , "GeoLocation" :
    {
      "Latitude": 40.75273
      , "Longitude": -74.006408
      , "Location": "(40.75273, -74.006408)"
    }
  , "Type of Base": "BLACK CAR BASE"
  , "Date": "08/15/2019"
  , "Time": "18:03:31"
}

```

### Sample Data: TaxiZones.csv

LocationID	1
Borough	EWR
Zone	Newark Airport
service_zone	EWR

### Sample Data: green\_tripdata\_2019-05.csv/ green\_tripdata\_2019-06.csv

VendorID	1
lpep_pickup_datetime	2019-05-01 00:48:55
lpep_dropoff_datetime	2019-05-01 00:55:07
store_and_fwd_flag	N
RatecodeID	1
PULocationID	41
DOLocationID	42
passenger_count	1
trip_distance	1.50
fare_amount	7.5
extra	0
tip_amount	0.5
mta_tax	0
tolls_amount	0
ehail_fee	Null
improvement_surcharge	0.3
total_amount	8.3
payment_type	2
trip_type	1
congestion_surcharge	0

### Sample Data: PaymentsType.json







```
{ "PaymentTypeID": 1, "PaymentType": "Credit Card" }
{ "PaymentTypeID": 2, "PaymentType": "Cash" }
{ "PaymentTypeID": 3, "PaymentType": "No Charge" }
{ "PaymentTypeID": 4, "PaymentType": "Dispute" }
{ "PaymentTypeID": 5, "PaymentType": "Unknown" }
{ "PaymentTypeID": 6, "PaymentType": "Voided Trip" }
```




### Sample Data: RateCodes.csv





RateCodeID	1
RateCode	Standard Rate
IsApproved	Yes






## Source Data Files







### New York Taxi datasets

<input type="checkbox"/> Name
 common
 fhv
 green
 yellow
 NycTaxi-Metadata
 nyctaxi-metadata

Windows (C:) > nyctaxi-datasets > common				
<input type="checkbox"/> Name	Date modified	Type	Size	
 PaymentTypes	7/11/2020 7:17 PM	JSON Source File	1 KB	
 RateCodes	7/11/2020 7:16 PM	Microsoft Excel Comma Separated Values File	1 KB	
 TaxiZones	7/11/2020 7:16 PM	Microsoft Excel Comma Separated Values File	13 KB	

Windows (C:) > nyctaxi-datasets > yellow				
<input type="checkbox"/> Name	Date modified	Type	Size	
 yellow_tripdata_2019_01	4/12/2020 8:41 PM	Microsoft Excel Comma Separated Values File	1,311 KB	
 yellow_tripdata_2019_02	4/12/2020 8:38 PM	Microsoft Excel Comma Separated Values File	1,325 KB	
 yellow_tripdata_2019_05	7/11/2020 7:12 PM	Microsoft Excel Comma Separated Values File	4,526 KB	
 yellow_tripdata_2019_06	7/11/2020 7:14 PM	Microsoft Excel Comma Separated Values File	4,520 KB	

Windows (C:) > nyctaxi-datasets > green				
<input type="checkbox"/> Name	Date modified	Type	Size	
 green_tripdata_2019_01	4/10/2020 9:02 PM	Microsoft Excel Comma Separated Values File	1,754 KB	
 green_tripdata_2019_02	4/10/2020 9:00 PM	Microsoft Excel Comma Separated Values File	1,320 KB	
 green_tripdata_2019-05	7/11/2020 7:10 PM	Microsoft Excel Comma Separated Values File	4,490 KB	
 green_tripdata_2019-06	7/11/2020 7:11 PM	Microsoft Excel Comma Separated Values File	4,495 KB	
 green_tripdata_json_2019-02	4/21/2020 3:20 PM	JSON Source File	6,579 KB	

Windows (C:) > nyctaxi-datasets > fhv				
<input type="checkbox"/> Name	Date modified	Type	Size	
 fhv_bases_extra	2/7/2020 4:11 PM	Microsoft Excel Comma Separated Values File	150 KB	
 fhv_tripdata_2019_01	4/14/2020 8:31 PM	Microsoft Excel Comma Separated Values File	807 KB	
 fhv_tripdata_2019_02	4/14/2020 8:30 PM	Microsoft Excel Comma Separated Values File	820 KB	
 fhv_tripdata_2019-05	7/11/2020 7:18 PM	Microsoft Excel Comma Separated Values File	2,729 KB	
 fhv_tripdata_2019-06	7/11/2020 7:19 PM	Microsoft Excel Comma Separated Values File	2,730 KB	
 FhvBases	2/27/2020 5:22 PM	JSON Source File	437 KB	

# NycTaxiMetaData

## NYC\_TAXI DATA

### RateCodes Data

RateCodeID	RateCodeID,RateCode,IsApproved
RateCode	1,Standard Rate,Yes
IsApproved	

green_tripdata	Sample Data	yellow_tripdata	Sample Data
VendorID	2	VendorID	1
lpep_pickup_datetime	2019-06-01 00:25:27	lpep_pickup_datetime	6/1/2019 0:55
lpep_dropoff_datetime	2019-06-01 00:33:52	lpep_dropoff_datetime	6/1/2019 0:56
store and fwd flag	N	passenger count	1
RateCodeID	1	trip distance	0
PULocationID	74	RateCodeID	1
DOLocationID	263	store and fwd flag	N
passenger count	5	PULocationID	145
trip distance	2.34	DOLocationID	145
fare amount	9	payment_type	2
Extra	0.5	fare amount	3
mta tax	0.5	Extra	0.5
tip amount	1	mta tax	0.5
tolls amount	0	tip amount	0
shail fee		tolls amount	0
improvement surcharge	0.3	improvement surcharge	0.3
total amount	14.05	total amount	4.3
payment_type	1	congestion surcharge	0
trip_type	1		
congestion surcharge	2.75		

### PaymentTypes Data

PaymentTypeID
PaymentType

```
{ "PaymentTypeID":1,
  "PaymentType":"Credit Card" }
```

### TaxiZones Data

LocationID
Borough
Zone
service zone

3

# NycTaxiMetaData

## NYC\_TAXI DATA

### fhv\_tripdata

dispatching_base_num (FK)
pickup_datetime
dropoff_datetime
PULocationID (FK)
DOLocationID (FK)
SR_Flag

### fhv\_basedata

License Number (PK)					
Entity Name					
Telephone Number					
SHL Endorsed					
Address	Building	Street	City	State	Postcode
GeoLocation	Latitude	Longitude	Location		
Type of Base					
Date					
Time					

### TaxiZones Data

LocationID (PK)
Borough
Zone
service zone

4

## New York Taxi Data Sets Metadata

### Green Trip Data Sets Metadata with Data Type ( 20 columns)

<i>Column Name</i>	<i>Data type</i>
VendorID	integer
lpep_pickup_datetime	timestamp
lpep_dropoff_datetime	timestamp
store_and_fwd_flag	string
RatecodeID	integer
PULocationID	integer
DOLocationID	integer
passenger_count	integer
trip_distance	double
fare_amount	double
extra	double
mta_tax	double
tip_amount	double
tolls_amount	double
ehail_fee	string
improvement_surcharge	double
total_amount	double
payment_type	integer
trip_type	integer
congestion_surcharge	double

### Yellow Trip Data Sets Metadata with Data Type ( 18 columns)

<i>Column Name</i>	<i>Data type</i>
VendorID	integer
tppep_pickup_datetime	timestamp
tppep_dropoff_datetime	timestamp
passenger_count	integer
trip_distance	double
RatecodeID	integer
store_and_fwd_flag	string
PULocationID	integer
DOLocationID	integer
payment_type	integer
fare_amount	double
extra	double
mta_tax	double
tip_amount	double
tolls_amount	double
improvement_surcharge	double
total_amount	double
congestion_surcharge	double

### **Taxi Zones Data Set Metadata with Data Type ( 4 columns)**

<b>Column Name</b>	<b>Data type</b>
LocationID	integer
Borough	string
Zone	string
service_zone	string

### **Payment Types Data Set Metadata with Data Type ( 2 columns)**

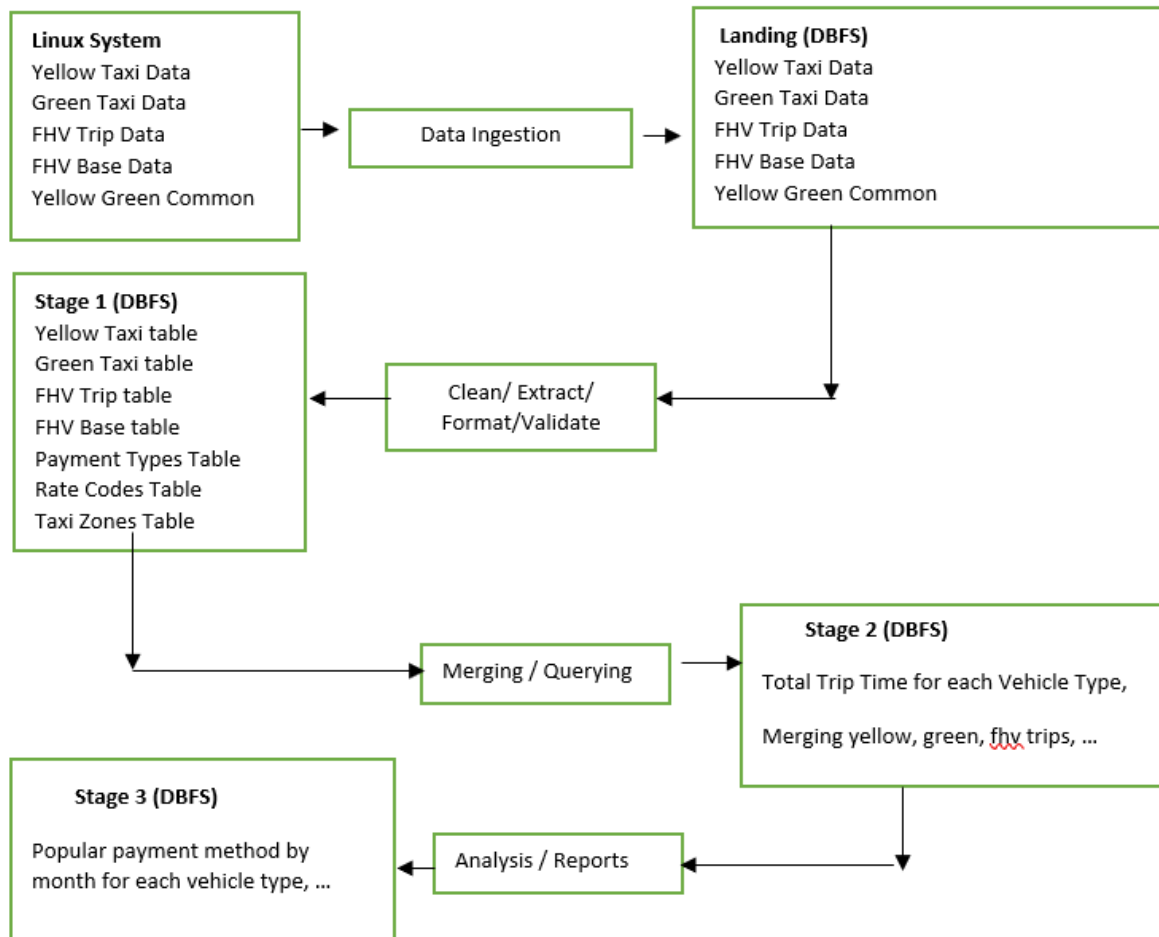
<b>Column Name</b>	<b>Data type</b>
PaymentTypeID	integer
PaymentType	string

### **Rate Codes Data Set Metadata with Data Type ( 3 columns)**

<b>Column Name</b>	<b>Data type</b>
RateCodeID	integer
RateCode	string
IsApproved	string

# New York Taxi Case Study Dataflow Diagram

## Dataflow Diagram:





# Path Convention for New York Taxi Data sets in DBFS

## (For batch data processing)

dbfs:/FileStore/tables/<CUSTOMIZED LOCATION>

### Landing location of Source Data in DBFS

dbfs:/FileStore/tables/nyctaxidata/landing/greendata/green\_tripdata  
dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow\_tripdata  
dbfs:/FileStore/tables/nyctaxidata/landing/yellow\_green\_commondata/ratecodes/  
dbfs:/FileStore/tables/nyctaxidata/landing/yellow\_green\_commondata/paymenttypes/  
dbfs:/FileStore/tables/nyctaxidata/landing/yellow\_green\_commondata/taxizones/

dbfs:/FileStore/tables/nyctaxidata/landing/fhvdata/fhv\_tripdata  
dbfs:/FileStore/tables/nyctaxidata/landing/fhvdata/fhvbases  
dbfs:/FileStore/tables/nyctaxidata/landing/fhvdata/fhvbases\_csv

### Stage1 Location of Data after ETL/Cleaning Operations from Landing Location:

dbfs:/FileStore/tables/nyctaxidata/stage1\_spark/greendata/green\_tripdata  
dbfs:/FileStore/tables/nyctaxidata/stage1\_spark/yellowdata/yellow\_tripdata  
dbfs:/FileStore/tables/nyctaxidata/stage1\_spark/yellow\_green\_commondata/ratecodes  
dbfs:/FileStore/tables/nyctaxidata/stage1\_spark/yellow\_green\_commondata/paymenttypes  
dbfs:/FileStore/tables/nyctaxidata/stage1\_spark/yellow\_green\_commondata/taxizones

dbfs:/FileStore/tables/nyctaxidata/stage1\_spark/fhvdata/fhv\_tripdata  
dbfs:/FileStore/tables/nyctaxidata/stage1\_spark/fhvdata/fhv\_bases  
dbfs:/FileStore/tables/nyctaxidata/stage1\_spark/fhvdata/taxizones

### Merged Table location (Stage2)

dbfs:/FileStore/tables/nyctaxidata/stage2\_spark/

### Processed/Output Files locations (Stage3)

dbfs:/FileStore/tables/nyctaxidata/stage3\_spark/

### FILE LOCATION FOR STRUCTURED STREAMING

dbfs:/FileStore/tables/nyctaxidata/streaming/yellowdata/yellow\_tripdata

# COMMANDS TO WORK WITH FILES IN DBFS

<https://docs.databricks.com/dbfs/index.html>

<https://docs.databricks.com/files/index.html>

## Command to check the Contents of file in DBFS

```
dbutils.fs.head("dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_05.csv ")
```

```
%fs head
dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_05.csv
```

## Command to remove file from DBFS

```
dbutils.fs.rm("/FileStore/tables/your_table_name.csv")
```

## Command to check DBFS root

### # Default location for %fs is root

```
%fs ls /tmp/
%fs mkdirs /tmp/my_cloud_dir
%fs cp /tmp/test_dbfs.txt /tmp/file_b.txt
```

### # Default location for dbutils.fs is root

```
dbutils.fs.ls ("/tmp/")
dbutils.fs.put("/tmp/my_new_file", "This is a file in cloud storage.")
```

### # Default location for %sh is the local filesystem

```
-----
%sh ls /dbfs/tmp/
```

### # Default location for os commands is the local filesystem

```
-----
import os
os.listdir('/dbfs/tmp')
```

# With **%fs** and **dbutils.fs**, you must use **file:/** to read from local filesystem

---

```
%fs ls file:/tmp
```

```
%fs mkdirs file:/tmp/my_local_dir
```

```
dbutils.fs.ls ("file:/tmp/")
```

```
dbutils.fs.put("file:/tmp/my_new_file", "This is a file on the  
local driver node.")
```

**# %sh reads from the local filesystem by default**

---

```
%sh ls /tmp
```

# Distributed Data processing using Spark Core (RDD)

## Working with yellowtrip data set

**Sample Data: yellow\_tripdata\_2019-05.csv/ yellow\_tripdata\_2019-06.csv**

VendorID	1	attribute[0]
tpep_pickup_datetime	2019-05-01 00:14:50	attribute[1]
tpep_dropoff_datetime	2019-05-01 00:16:48	attribute[2]
passenger_count	1	attribute[3]
trip_distance	.00	attribute[4]
RatecodeID	1	attribute[5]
store_and_fwd_flag	N	attribute[6]
PULocationID	145	attribute[7]
DOLocationID	145	attribute[8]
payment_type	2	attribute[9]
fare_amount	3	attribute[10]
extra	0.5	attribute[11]
mta_tax	0.5	attribute[12]
tip_amount	0	attribute[13]
tolls_amount	0	attribute[14]
improvement_surcharge	0.3	attribute[15]
total_amount	4.3	attribute[16]
congestion_surcharge	0	attribute[17]

## Problem Statements:-

(Develop Notebook applications using Spark RDD API to generate reports as per the problem statements)

### 1. Ingest All Yellow Trip data Files into DBFS Landing location

***/FileStore/tables/nyctaxidata/landing/yellowdata/yellow\_tripdata***

### 2. Create RDD to store valid yellow trip records for June Month of Year 2019

(Filter conditions to be checked)

- tpep\_pickup\_datetime and tpep\_dropoff\_datetime should not be same
- passenger\_count should not be zero
- trip\_distance should not be zero
- fare\_amount should not be zero
- total\_amount should not be zero

## Filtering Yellow Trip Data sets

Sample Data: yellow\_tripdata\_2019-05.csv/ yellow\_tripdata\_2019-06.csv

VendorID	1	attribute[0]
tpep_pickup_datetime	2019-05-01 00:14:50	attribute[1]
tpep_dropoff_datetime	2019-05-01 00:16:48	attribute[2]
passenger_count	1	attribute[3]
trip_distance	.00	attribute[4]
RatecodeID	1	attribute[5]
store_and_fwd_flag	N	attribute[6]
PULocationID	145	attribute[7]
DOLocationID	145	attribute[8]
payment_type	2	attribute[9]
fare_amount	3	attribute[10]
extra	0.5	attribute[11]
mta_tax	0.5	attribute[12]
tip_amount	0	attribute[13]
tolls_amount	0	attribute[14]
improvement_surcharge	0.3	attribute[15]
total_amount	4.3	attribute[16]
congestion_surcharge	0	attribute[17]

3. Create an application to get the total valid trip count of yellow taxi for each Vendor for June month
4. Create an application to find total trip cost for each yellow taxi Vendor for June month for all valid trips
5. Create an application to get the report of total passenger count for each yellow taxi Vendor for June month for all valid trips
6. Create an application to generate the report of total trip distance for each yellow taxi Vendor for June month for all valid trips
7. Create an application to generate the report of **total passenger count for each day for June month (Consider pickup\_datetime)** for all valid trips

VendorID,tpep\_pickup\_datetime,tpep\_dropoff\_datetime,passenger\_count,trip\_distance,Ratecode ID,store\_and\_fwd\_flag,PULocationID,DOLocationID,payment\_type,fare\_amount,extra,mta\_tax,tip\_amount,tolls\_amount,improvement\_surcharge,total\_amount,congestion\_surcharge

1,2019-01-01 00:46:40,2019-01-01 00:53:20,1,1.50,1,N,151,239,1,7,0.5,0.5,1.65,0,0.3,9.95,

8. Create an application to generate the report of total passenger count for each Vendor and each day of June month (Consider pickup\_datetime) for all valid trips

9. Save the previous results into DBFS location

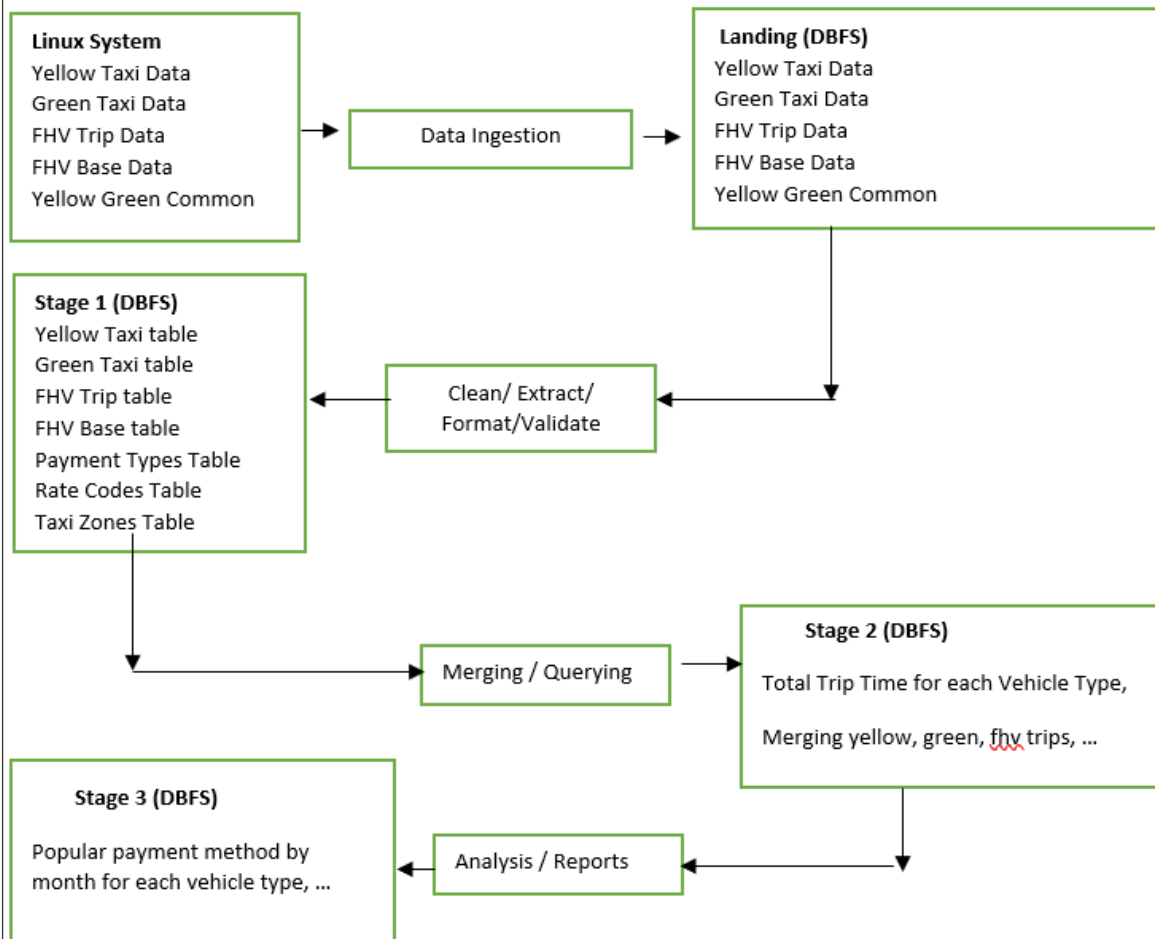
```
dbfs:/FileStore/tables/nyctaxidata/stage3_spark/passcount_vendorday
```

10. Check the DBFS location

---

# Distributed Data Processing Using Spark SQL

## Dataflow Diagram:



# Data Ingestion and Data Preparation

Create separate notebook for each problem statement

## 1. Ingest all New York Taxi datasets into DBFS Landing locations and check the locations

- a. Yellow Trip data sets
- b. Green Trip data sets
- c. Taxi zones data sets
- d. Payment types data sets
- e. Rate codes data sets

## 2. ETL Processing with Yellow Trip Data Sets

- a. Create a temporary view "**YellowTripRawDataView**" from Landing location
- b. Filter the YellowTripRawDataView data as per the below filter conditions –
  1. tpep\_pickup\_datetime and tpep\_pickup\_datetime should not be same
  2. record will be considered on for year 2019 only
  3. passenger\_count should not be zero
  4. trip\_distance should not be zero
  5. fare\_amount should not be zero
  6. total\_amount should not be zero
  7. PULocationID should not be null
  8. DOLocationID should not be null
- c. Create a Delta Table "**YellowTripDelta**" from the filtered view with the listed columns –  
( "VendorID", "PickupTime", "DropTime", "PassengerCount", "TripDistance", "RatecodeID",  
"PickupLocationId", "DropLocationId", "payment\_type", "fare\_amount", "extra",  
"mta\_tax", "tip\_amount", "tolls\_amount", "improvement\_surcharge", "total\_amount",  
"congestion\_surcharge")  
# Ignore the store\_and\_fwd\_flag column

## 3. ETL Processing with Green Trip Data Sets

- a. Create a temporary view "**GreenTripRawDataView**" from Landing location
- b. Filter the GreenTripRawDataView data as per the below filter conditions –
  1. lpep\_pickup\_datetime and lpep\_pickup\_datetime should not be same
  2. record will be considered on for year 2019
  3. passenger\_count should not be zero
  4. trip\_distance should not be zero
  5. fare\_amount should not be zero



6. total\_amount should not be zero
7. PULocationID should not be null
8. DOLocationID should not be null

- c. Create a Delta Table **"GreenTripDelta"** from the filtered view with the listed columns –

```
( "VendorID","PickupTime","DropTime","PassengerCount","TripDistance","RatecodeID",
  "PickupLocationID","DropLocationID","payment_type", "fare_amount", "extra",
  "mta_tax","tip_amount","tolls_amount", "improvement_surcharge", "total_amount",
  "congestion_surcharge", "ehail_fee", "trip_type")
# Ignore the store_and_fwd_flag column
```

#### 4. Combine Yellow and Green Trip Data Sets

- a. Create a Delta Table **"YellowGreenTripCombineDelta"** combining **YellowTripDelta** and **GreenTripDelta** tables with listed columns –

```
("VendorID","PickupTime","DropTime","PassengerCount","TripDistance","RatecodeID","PickupL
ocationID","DropLocationID","payment_type",
"fare_amount","extra","mta_tax","tip_amount","tolls_amount",
"improvement_surcharge","total_amount","congestion_surcharge","taxiType")
```

# Where **taxiType** column value will be **"Green"** for all green trip records and **"Yellow"** for all yellow trip records.

#### 5. Working with TaxiZones data sets

- a. Create a temporary view **"TaxiZonesRawDataView"** from Landing location
- b. Create a Delta Table **"TaxiZonesDelta"** from the **TaxiZonesRawDataView** view with all columns

#### 6. Working with PaymentTypes data sets

- a. Create a temporary view **"PaymentTypesRawDataView"** from Landing location
- b. Create a Delta Table **"PaymentTypesDelta"** from the **PaymentTypesRawDataView** view with all columns

#### 7. Working with RateCodes data sets

- a. Create a temporary view **"RateCodesRawDataView"** from Landing location
- b. Create a Delta Table **"RateCodesDelta"** from the **RateCodesRawDataView** view with all columns

## Create Report on Yellow Green Combined Data sets

### (Use Single Notebook for all reports)

1. Generate Report to get total trip time in hours for each taxi type
2. Generate Report to get taxi-type-wise total trip time in hours for each trip-month
3. Generate Report to get taxi-type-wise total number of passengers for each trip month
4. Generate Report to get taxi-type-wise total number of payments for each payment-type
5. Generate Report to get the total number of light trips for each taxi type (when passenger count for each trip is  $\leq 2$ )
6. Generate Report to get the total number of light trips for each taxi type month-wise (when passenger count for each trip is  $\leq 2$ ) and save the results in a delta table  
"lightTripsTaxiTypeMonWise"
7. Generate Report to get the total number of fully-loaded trips for each taxi type (when passenger count for each trip is  $\geq 4$ )
8. Generate Report to get the total number of fully-loaded trips for each taxi type month-wise (when passenger count for each trip is  $\geq 4$ ) and save the results in a delta table  
"loadedTripsTaxiTypeMonWise"
9. Generate Report to get the total number of midnight trips for each taxi type (when trips happen between 12AM to 4AM)
10. Generate Report to get the total number of midnight trips for each taxi type month-wise (when trips happen between 12AM to 4AM)



# Create Report on Yellow Green Combined & Common Data sets

## (Use Single Notebook for all reports)

1. Create a Report to get total trip time in hours for each taxi type
2. Create a Report to generate Passenger Count for Each Zone  
(Join **YellowGreenTripCombineDelta** and **TaxiZonesDelta** tables for creating report)  
Display and Save the output in table "**ZonePassCountTable**".
3. Create a database "**nyctaxi**" in Spark SQL
4. Create a Report to generate Total Trip Count Per Zone, TaxiType, TripMonth, VendorID .  
Display and save the output in table "**ZoneTaxiMonVendorTripCountTable**" under "**nyctaxi**" database.
5. Create a Report to generate Total Trip Time Per Borough,TaxiType,TripMonth,VendorID.  
Display and save the output in table "**BoroughTaxiMonVendorTripCountTable**" under "**nyctaxi**" database.
6. Create a Report to generate Total Travel Fare Per Service\_Zone,TaxiType,TripMonth. Display and Save the output in table "**TotalFareSZoneTaxiMonth**" under "**nyctaxi**" database.
7. Generate report on Total Different Payment Method count for Each taxi type, each tripMonth, each VendorID.  
(Join **YellowGreenTripCombineDelta** and **PaymentTypesDelta** tables for creating report)  
Display and Save the output in table "**PaymentCountTaxiMonthVendor**" under "**nyctaxi**" database.

Create a Report to generate Total Different Payment Method count for Each Zone, Each taxi type, each tripMonth,each VendorID. Display and Save the output in table "**PaymentCountZoneTaxiMonthVendor**" under "**nyctaxi**" database.

8. Create a report to generate Total Trip count for each vendor for each Zone where "Standard Rate" has been applied. Display and Save the output in table "**StandardRateTripCount**" under "**nyctaxi**" database.
9. Check "**nyctaxi**" database and display the records from each table.

## Incremental Data Ingestion with Auto Loader

Incremental ETL is important since it allows us to deal solely with new data that has been encountered since the last ingestion.

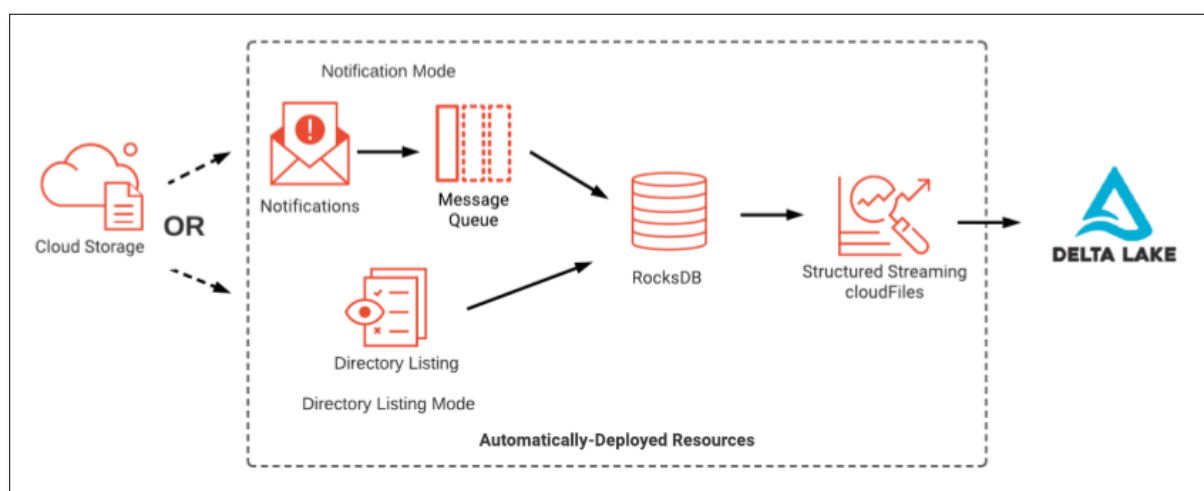
Reliably processing only the new data reduces redundant processing and helps enterprises reliably scale data pipelines.

The first step for any successful data lakehouse implementation is ingesting into a Delta Lake table from cloud storage.

Historically, ingesting files from a data lake into a database has been a complicated process.

Databricks Auto Loader provides an easy-to-use mechanism for incrementally and efficiently processing new data files as they arrive in cloud file storage.

In this tutorial, you'll see Auto Loader in action.



### Learning Objectives

By the end of this lesson, you should be able to:

- Execute Auto Loader code to incrementally ingest data from cloud storage to Delta Lake
- Describe what happens when a new file arrives in a directory configured for Auto Loader
- Query a table fed by a streaming Auto Loader query

## Dataset Used

This demo uses **yellow trip** datasets from NycTaxi datasets which is in CSV format.

### Column List

VendorID, tpep\_pickup\_datetime, tpep\_dropoff\_datetime, passenger\_count, trip\_distance, RatecodeID, store\_and\_fwd\_flag, PULocationID, DOLocationID, payment\_type, fare\_amount, extra, mta\_tax, tip\_amount, tolls\_amount, improvement\_surcharge, total\_amount, congestion\_surcharge

## Getting Started

**Note :- All commands to be issued in Databricks Notebook**

Run the following cell to check **Yellow** trip data files

```
%fs ls dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata
```

Table				
	path	name	size	modificationTime
1	dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_01.csv	yellow_tripdata_2019_01.csv	1341900	1668675324000
2	dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_05.csv	yellow_tripdata_2019_05.csv	4633703	1668673115000
3	dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_06.csv	yellow_tripdata_2019_06.csv	4628092	1668673569000

Showing all 3 rows.

## Delete all the files from the location

```
%fs rm dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_01.csv
```

```
%fs rm dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_01.csv
```

```
res6: Boolean = true
```

```
%fs rm dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_06.csv
```

```
%fs rm dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_06.csv
```

```
res7: Boolean = true
```

## Load **yellow\_tripdata\_2019\_05.csv** File in DBFS as per the Path Convention

dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow\_tripdata

### Check the file in DBFS

%fs ls dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow\_tripdata

Table				
	path	name	size	modificationTime
1	dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_05.csv	yellow_tripdata_2019_05.csv	4633703	1668673115000

Showing 1 row.

## Using Auto Loader

In the cell below, a function is defined to demonstrate using Databricks Auto Loader with the PySpark API. This code includes both a **Structured Streaming read and write**.

Note that when using Auto Loader with **automatic schema inference and evolution**, the 4 arguments shown here should allow ingestion of most datasets.

These arguments are explained below.

argument	what it is	how it's used
<b>data_source</b>	The directory of the source data	Auto Loader will detect new files as they arrive in this location and queue them for ingestion; passed to the <b>.load()</b> method
<b>source_format</b>	The format of the source data	While the format for all Auto Loader queries will be <b>cloudFiles</b> , the format of the source data should always be specified for the <b>cloudFiles.format</b> option
<b>table_name</b>	The name of the target table	Spark Structured Streaming supports writing directly to Delta Lake tables by passing a table name as a string to the <b>.table()</b> method. Note that you can either append to an existing table or create a new table
<b>checkpoint_directory</b>	The location for storing metadata about the stream	This argument is passed to the <b>checkpointLocation</b> and <b>cloudFiles.schemaLocation</b> options. Checkpoints keep track of streaming progress, while the schema location tracks updates to the fields in the source dataset

The code below has been streamlined to demonstrate Auto Loader functionality.

```
def autoloader_to_table(data_source, source_format, table_name, checkpoint_directory):  
    query = (spark.readStream  
              .format("cloudFiles")  
              .option("cloudFiles.format", source_format)  
              .option("cloudFiles.schemaLocation", checkpoint_directory)  
              .load(data_source)  
              .writeStream  
              .option("checkpointLocation", checkpoint_directory)  
              .option("mergeSchema", "true")  
              .table(table_name))  
    return query
```

In the following cell, we use the previously defined function to begin an Auto Loader stream.

Here, we're reading from a source directory of CSV files.

```
query = autoloader_to_table(data_source =  
    "dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata",  
    source_format = "csv",  
    table_name = "yellowtripsrc",  
    checkpoint_directory =  
    "dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/checkpoint/yellowtripsrc")
```

```
query = autoloader_to_table(data_source = "dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata",  
    source_format = "csv",  
    table_name = "yellowtripsrc",  
    checkpoint_directory = "dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/checkpoint/yellowtripsrc")
```

8d55517f-f774-4a83-bf81-40bc3b5c7668 Last updated: 16 days ago

Because Auto Loader uses Spark Structured Streaming to load data incrementally, the code above **doesn't appear to finish executing**.

We can think of this as a **continuously active query**. This means that as soon as new data arrives in our data source, it will be processed through our logic and loaded into our target table.



## Query the Target Table Data

```
%sql
SELECT * FROM yellowtripsrc
```

Table											
	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	passenger_count	trip_distance	RatecodeID	store_and_fwd_flag	PULocationID	DOLocationID	payment_type	fare
1	1	2019-05-01 00:14:50	2019-05-01 00:16:48	1	.00	1	N	145	145	2	3
2	1	2019-05-01 00:35:54	2019-05-01 00:37:27	1	1.50	1	N	145	145	2	3
3	1	2019-05-01 00:37:45	2019-05-01 00:37:49	1	1.50	1	N	145	145	2	2.5
4	1	2019-05-01 00:44:57	2019-05-01 00:50:11	1	.70	1	N	161	161	2	5
5	1	2019-05-01 00:59:48	2019-05-01 01:10:22	1	2.00	1	N	163	141	1	9.5
6	2	2019-05-01 00:48:04	2019-05-01 00:49:00	1	.00	1	N	193	193	2	2.5
7	2	2019-05-01 00:05:26	2019-05-01 00:05:34	1	.00	1	N	193	193	2	2.5

Truncated results, showing first 1,000 rows.

```
%sql
DESCRIBE TABLE yellowtripsrc
```

Table			
	col_name	data_type	comment
1	VendorID	string	null
2	tpep_pickup_datetime	string	null
3	tpep_dropoff_datetime	string	null
4	passenger_count	string	null
5	trip_distance	string	null
6	RatecodeID	string	null
7	store and fwd flag	string	null

Showing all 19 rows.

Use the cell below to define a temporary view that summarizes the recordings in our target table.

We'll use this view below to demonstrate how new data is automatically ingested with Auto Loader.

```
%sql
CREATE OR REPLACE TEMP VIEW vendor_trip_count AS
SELECT VendorID, count(*) total_trip_count
FROM yellowtripsrc
GROUP BY VendorID;

SELECT * FROM vendor_trip_count
```

	VendorID ▲	total_trip_count ▲
1	1	18692
2	4	160
3	2	31147

Showing all 3 rows.

## Land New Data

Load another Yellow trip data file **yellow\_tripdata\_2019\_06.csv** in the path

**dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow\_tripdata**

As mentioned previously, Auto Loader is configured to incrementally process files from a directory in cloud object storage into a Delta Lake table.

We have configured and are currently executing a query to process CSV files from the location specified by **source\_path** into a table named **target\_table**.

Let's review the contents of the **source\_path** directory.

List the contents of the **source\_path** again using the cell below. You should see an additional CSV file .

```
%fs ls dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata
```

	path ▲	name ▲	size ▲	modificationTime ▲
1	dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_05.csv	yellow_tripdata_2019_05.csv	4633703	1668673115000
2	dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_06.csv	yellow_tripdata_2019_06.csv	4628092	1668673569000

Showing all 2 rows.

```
files =
dbutils.fs.ls(f"dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata
/yellow_tripdata")

display(files)
```

	path ▲	name ▲	size ▲	modificationTime ▲
1	dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_05.csv	yellow_tripdata_2019_05.csv	4633703	1668673115000
2	dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_06.csv	yellow_tripdata_2019_06.csv	4628092	1668673569000

Showing all 2 rows.

## Tracking Ingestion Progress

Historically, many systems have been configured to either reprocess all records in a source directory to calculate current results or require data engineers to implement custom logic to identify new data that's arrived since the last time a table was updated.

With Auto Loader, your table has already been updated.

Run the query below to confirm that new data has been ingested.

```
%sql  
SELECT * FROM vendor_trip_count
```

Table		
	VendorID	total_trip_count
1	1	36802
2	4	328
3	2	62868

Showing all 3 rows.

The Auto Loader query we configured earlier automatically detects and processes records from the source directory into the target table.

There is a slight delay as records are ingested, but an Auto Loader query executing with default streaming configuration should update results in near real time.

The query below shows the table history.

A new table version should be indicated for each **STREAMING UPDATE**.  
These update events coincide with new batches of data arriving at the source.

```
%sql  
DESCRIBE HISTORY yellowtripsrc
```

Table						
	version	timestamp	userId	userName	operation	operationParameters
1	3	2022-11-17T08:26:15.000+0000	8780831625071985	raju.chal@accenture.com	STREAMING UPDATE	{ "outputMode": "Append", "queryId": "8d55517f-f774-4a83-bf81-40bc3b5c7668", "epochId": "2" }
2	2	2022-11-17T08:26:10.000+0000	8780831625071985	raju.chal@accenture.com	STREAMING UPDATE	{ "outputMode": "Append", "queryId": "8d55517f-f774-4a83-bf81-40bc3b5c7668", "epochId": "1" }
3	1	2022-11-17T08:21:40.000+0000	8780831625071985	raju.chal@accenture.com	STREAMING UPDATE	{ "outputMode": "Append", "queryId": "8d55517f-f774-4a83-bf81-40bc3b5c7668", "epochId": "0" }
4	0	2022-11-17T08:21:20.000+0000	8780831625071985	raju.chal@accenture.com	CREATE TABLE	{ "isManaged": "true", "description": null, "partitionBy": "[]", "properties": "{}" }

Showing all 4 rows.

## Create a Streaming View

```
(spark.readStream
  .table("yellowtripsrc")
  .createOrReplaceTempView("streaming_yellowtrip_vw"))
```

```
%sql
SELECT * FROM streaming_yellowtrip_vw
```

%sql  
SELECT \* FROM streaming\_yellowtrip\_vw

display\_query\_1 (id: 5a856a6d-963d-40b7-b2e5-3fc1391f5d83) Last updated: 16 days ago

Table

	VendorID	tpcp_pickup_datetime	tpcp_dropoff_datetime	passenger_count	trip_distance	RatecodeID	store_and_fwd_flag	PULocationID	DOLocationID	payment_type	fare_amount
1	1	2019-06-01 00:55:13	2019-06-01 00:56:17	1	.00	1	N	145	145	2	3
2	1	2019-06-01 00:06:31	2019-06-01 00:06:52	1	.00	1	N	262	263	2	2.5
3	1	2019-06-01 00:17:05	2019-06-01 00:36:38	1	4.40	1	N	74	7	2	17.5
4	1	2019-06-01 00:59:02	2019-06-01 00:59:12	0	.80	1	N	145	145	2	2.5
5	1	2019-06-01 00:03:25	2019-06-01 00:15:42	1	1.70	1	N	113	148	1	9.5
6	1	2019-06-01 00:28:31	2019-06-01 00:39:23	2	1.60	1	N	79	125	1	9.5
7	1	2019-06-01 00:46:46	2019-06-01 00:50:55	4	.60	1	N	211	148	2	4.5

Truncated results, showing first 1,000 rows.

## Cancel the streaming Execution

## Find the Total Trip Count for Each Vendor ID

```
%sql
SELECT VendorID, count(*) total_trip_count
FROM streaming_yellowtrip_vw
GROUP BY VendorID;
```

%sql  
SELECT VendorID, count(\*) total\_trip\_count  
FROM streaming\_yellowtrip\_vw  
GROUP BY VendorID;

display\_query\_2 (id: d89beb83-e95f-4285-8644-4be881454398) Last updated: 16 days ago

Table

	VendorID	total_trip_count
1	1	36802
2	4	328
3	2	62868

Showing all 3 rows.

## Create another Temp View to store the Total Trip Count for Each Vendor ID

```
%sql
CREATE OR REPLACE TEMP VIEW vendor_trip_count_vw AS
SELECT VendorID, count(*) total_trip_count
FROM streaming_yellowtrip_vw
GROUP BY VendorID;
```

## Write the TempView Records in another Table

```
(spark.table("vendor_trip_count_vw")
  .writeStream
  .option("checkpointLocation",
f"dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/checkpoint/vendor_trip_count_table")
  .outputMode("complete")
  .trigger(availableNow=True)
  .table("vendor_trip_count_table")
  .awaitTermination()
# This optional method blocks execution of the next cell until the incremental batch write has
succeeded
)
```

## Check the records from the Table

```
%sql
SELECT *
FROM vendor_trip_count_table;
```

Table		
	VendorID ▲	total_trip_count ▲
1	1	36802
2	4	328
3	2	62868

Showing all 3 rows.

## Write the TempView Records in another Table at some time interval

```
query = (spark.table("vendor_trip_count_vw")
  .writeStream
  .option("checkpointLocation",
f"dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/checkpoint/vendor_trip_count_table")
  .outputMode("complete")
  .trigger(processingTime='4 seconds')
  .table("vendor_trip_count_table"))
```

```
query = (spark.table("vendor_trip_count_vw")
        .writeStream
        .option("checkpointLocation", f"dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/checkpoint/vendor_trip_count_table")
        .outputMode("complete")
        .trigger(processingTime='4 seconds')
        .table("vendor_trip_count_table"))
```

▶ 0ba7b8ce-4b58-49a8-bdcb-ac7dcbe5e85 Last updated: 16 days ago

## Check the records from the Table

```
%sql
SELECT *
FROM vendor_trip_count_table
```

	VendorID ▲	total_trip_count ▲
1	1	36802
2	4	328
3	2	62868

Showing all 3 rows.

## Load Another New Yellow Trip Data File in DBFS

Load another Yellow trip data file **yellow\_tripdata\_2019\_01.csv** in the path

**dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow\_tripdata**

```
%sql
SELECT *
FROM vendor_trip_count_table
```

	VendorID ▲	total_trip_count ▲
1	1	42367
2	4	517
3	2	72114

Showing all 3 rows.

## **Do it Yourself :-**

**Implement Auto Loader application as per previous activities for Green Trip datasets to generate Total Passenger count for each Green Trip Vendor ID**

## Incremental Multi-Hop in the Lakehouse

Now that we have a better understanding of how to work with incremental data processing by combining Structured Streaming APIs and Spark SQL, we can explore the tight integration between **Structured Streaming and Delta Lake**.

### Learning Objectives

By the end of this lesson, you should be able to:

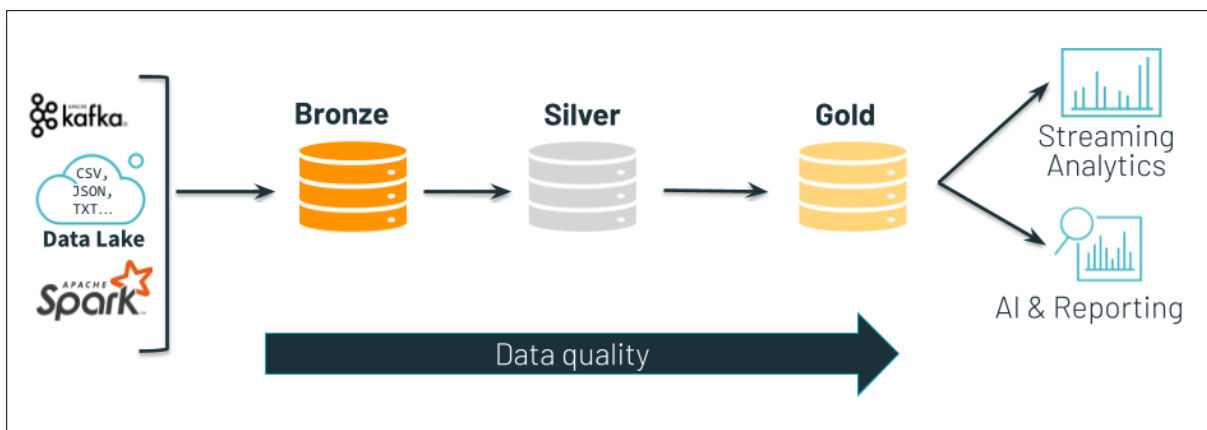
- Describe Bronze, Silver, and Gold tables
- Create a Delta Lake multi-hop pipeline

### Incremental Updates in the Lakehouse

Delta Lake allows users to easily combine streaming and batch workloads in a unified multi-hop pipeline.

Each stage of the pipeline represents a state of our data valuable to driving core use cases within the business.

Because all data and metadata lives in object storage in the cloud, multiple users and applications can access data in near-real time, allowing analysts to access the freshest data as it's being processed.



- **Bronze** tables contain raw data ingested from various sources (JSON files, RDBMS data, IoT data, to name a few examples).
- **Silver** tables provide a more refined view of our data. We can join fields from various bronze tables to enrich streaming records, or update account statuses based on recent activity.

- **Gold** tables provide business level aggregates often used for reporting and dashboarding. This would include aggregations such as daily active website users, weekly sales per store, or gross revenue per quarter by department.

The end outputs are actionable insights, dashboards and reports of business metrics.

By considering our business logic at all steps of the ETL pipeline, we can ensure that storage and compute costs are optimized by reducing unnecessary duplication of data and limiting ad hoc querying against full historic data.

**Each stage can be configured as a batch or streaming job, and ACID transactions ensure that we succeed or fail completely.**

## ***Datasets Used***

Use **yellow trip** datasets and **TaxiZones** datasets from NycTaxi datasets which is in CSV format.

The schema of our two datasets is represented below. Note that we will be manipulating these schema during various steps.

### **Column List of **yellow trip** datasets**

```
VendorID,tpep_pickup_datetime,tpep_dropoff_datetime,passenger_count,trip_
distance,RatecodeID,store_and_fwd_flag,PULocationID,DOLocationID,payment_
type,fare_amount,extra,mta_tax,tip_amount,tolls_amount,improvement_surcha
rge,total_amount,congestion_surcharge

1,2019-06-01 00:55:13,2019-06-01
00:56:17,1,.00,1,N,145,145,2,3,0.5,0.5,0,0,0.3,4.3,0
```

### **Column List of Taxizones datasets**

```
LocationID,Borough,Zone,service_zone

1,"EWR","Newark Airport","EWR"
```



## Getting Started

Run the following cell to check **Yellow** trip data files

```
%fs ls dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata
```

Table				
	path	name	size	modificationTime
1	dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_01.csv	yellow_tripdata_2019_01.csv	1341900	1668675324000
2	dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_05.csv	yellow_tripdata_2019_05.csv	4633703	1668673115000
3	dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_06.csv	yellow_tripdata_2019_06.csv	4628092	1668673569000

Showing all 3 rows.

## Delete two files from the location

```
%fs rm dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_01.csv
```

```
%fs rm dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_01.csv

res6: Boolean = true
```

```
%fs rm dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_06.csv
```

```
%fs rm dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_06.csv

res7: Boolean = true
```

## Check the location again

```
%fs ls dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata
```

Table				
	path	name	size	modificationTime
1	dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_05.csv	yellow_tripdata_2019_05.csv	4633703	1668673115000

Showing 1 row.

## Data Simulator

Databricks **Auto Loader** can automatically process files as they land in your cloud object stores.

To simulate this process, we will upload multiple Yellow trip data files in the same directory

`dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata`

### Bronze Table: Ingesting Raw JSON Recordings

Below, we configure a read on a raw CSV source using Auto Loader with schema inference.

Note that while you need to use the Spark DataFrame API to set up an incremental read, once configured you can immediately register a temp view to leverage Spark SQL for streaming transformations on your data.

#### NOTE:

For a CSV data source, Auto Loader will default to inferring each column as a **string**.

Here, we demonstrate specifying the data type for the `trip_distance` column using the `cloudFiles.schemaHints` option.

Note that specifying improper types for a field will result in null values.

```
(spark.readStream
  .format("cloudFiles")
  .option("cloudFiles.format", "csv")
  .option("cloudFiles.schemaHints", "trip_distance float")
  .option("cloudFiles.schemaLocation",
    f"dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/checkpoint/yellow_bronze")
  .load("dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata")
  .createOrReplaceTempView("yellowtrip_raw_temp"))
```

### Check the Schema of yellowtrip\_raw\_temp view

```
%sql
describe yellowtrip_raw_temp
```

Table				
	col_name	data_type	comment	
1	VendorID	string	null	
2	tpep_pickup_datetime	string	null	
3	tpep_dropoff_datetime	string	null	
4	passenger_count	string	null	
5	trip_distance	float	null	
6	RatecodeID	string	null	
7	store_and_fwd_flag	string	null	
Showing all 19 rows.				

## Enrich our RAW data sets

Here, we'll enrich our raw data with additional metadata describing the source file and the time it was ingested.

This additional metadata can be ignored during downstream processing while providing useful information for troubleshooting errors if corrupt data is encountered.

```
%sql
CREATE OR REPLACE TEMPORARY VIEW yellowtrip_bronze_temp AS (
  SELECT *, current_timestamp() receipt_time, input_file_name() source_file
  FROM yellowtrip_raw_temp
)
```

```
%sql
select * from yellowtrip_bronze_temp
```

Table											
	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	passenger_count	trip_distance	RatecodeID	store_and_fwd_flag	PULocationID	DOLocationID	payment_type	fare
1	1	2019-05-01 00:14:50	2019-05-01 00:16:48	1	0	1	N	145	145	2	3
2	1	2019-05-01 00:35:54	2019-05-01 00:37:27	1	1.5	1	N	145	145	2	3
3	1	2019-05-01 00:37:45	2019-05-01 00:37:49	1	1.5	1	N	145	145	2	2.5
4	1	2019-05-01 00:44:57	2019-05-01 00:50:11	1	0.7	1	N	161	161	2	5
5	1	2019-05-01 00:59:48	2019-05-01 01:10:22	1	2	1	N	163	141	1	9.5
6	2	2019-05-01 00:48:04	2019-05-01 00:49:00	1	0	1	N	193	193	2	2.5
7	2	2019-05-01 00:05:26	2019-05-01 00:05:34	1	0	1	N	193	193	2	2.5

The code below passes our enriched raw data back to PySpark API to process an incremental write to a Delta Lake table.

```
(spark.table("yellowtrip_bronze_temp")
  .writeStream
  .format("delta")
  .option("checkpointLocation",
f"dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/checkpoint/green_bronze")
  .outputMode("append")
  .table("yellowtrip_bronze"))
```

```
(spark.table("yellowtrip_bronze_temp")
  .writeStream
  .format("delta")
  .option("checkpointLocation", f"dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/checkpoint/green_bronze")
  .outputMode("append")
  .table("yellowtrip_bronze"))
```

▶  409c713a-8dcb-47a7-b08f-6766ff39b2bc Last updated: 32 days ago

Out[12]: <pyspark.sql.streaming.query.StreamingQuery at 0x7fca8a8f8760>

Check this stream will be **continuously** running

```
%sql
select count(*) from yellowtrip_bronze
```

```
%sql
select count(*) from yellowtrip_bronze
```

Table

	count(1)
1	49999

Showing 1 row.

## Upload another Yellow trip data files in the same directory

**dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow\_tripdata**

and you'll see the changes immediately detected by the streaming query you've written.

```
%fs ls dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata
```

```
%sql
select count(*) from yellowtrip_bronze
```

Table

	count(1)
1	99998

Showing 1 row.

## Load Static Lookup Table

The ACID guarantees that Delta Lake brings to your data are managed at the table level, ensuring that only fully successfully commits are reflected in your tables.

If you choose to merge these data with other data sources, be aware of how those sources version data and what sort of consistency guarantees they have.

In this simplified demo, we are loading a static CSV file (**TaxiZones** data file) to add TaxiZones data to our **YellowTrip** data sets.

In production, we could use Databricks' [Auto Loader](#) feature to keep an up-to-date view of these data in our Delta Lake.

### Ingest TaxiZones data sets into DBFS in this location

`/FileStore/tables/nyctaxidata/landing/yellow_green_commondata/taxizones/`

Check the file

```
%fs ls /FileStore/tables/nyctaxidata/landing/yellow_green_commondata/taxizones/
```

Table

	path	name	size	modificationTime
1	dbfs:/FileStore/tables/nyctaxidata/landing/yellow_green_commondata/taxizones/TaxiZones.csv	TaxiZones.csv	12320	1668678696000

Showing 1 row.

### Create View from TaxiZones data set

```
(spark.read
  .format("csv")
  .option("header", True)
  .option("inferSchema", True)
  .load(f'dbfs:/FileStore/tables/nyctaxidata/landing/yellow_green_commondata/taxizones/TaxiZones.csv')
  .createOrReplaceTempView("taxizones_view"))
```

```
%sql
SELECT * FROM taxizones_view
```

Table				
	LocationID	Borough	Zone	service_zone
1	1	EWB	Newark Airport	EWB
2	2	Queens	Jamaica Bay	Boro Zone
3	3	Bronx	Allerton/Pelham Gardens	Boro Zone
4	4	Manhattan	Alphabet City	Yellow Zone
5	5	Staten Island	Arden Heights	Boro Zone
6	6	Staten Island	Arrochar/Fort Wadsworth	Boro Zone
7	7	Queens	Astoria	Boro Zone

Showing all 265 rows.

```
%sql
describe taxizones_view
```

	col_name ▲	data_type ▲	comment ▲
1	LocationID	int	null
2	Borough	string	null
3	Zone	string	null
4	service_zone	string	null

Showing all 4 rows.

## Silver Table: Enriched Recording Data

As a second hop in our silver level, we will do the follow enrichments and checks:

- Our **YellowTrip** data will be joined with the **TaxiZones** to add Zone names
- tpep\_pickup\_datetime!= tpep\_dropoff\_datetime
- passenger\_count !=0
- trip\_distance!=0.0
- fare\_amount !=0.0
- total\_amount!=0.0

```
(spark.readStream
  .table("yellowtrip_bronze")
  .createOrReplaceTempView("yellowtrip_bronze_tmp"))
```

```
%sql
CREATE OR REPLACE TEMPORARY VIEW yellowtrip_taxizones_view AS (
  SELECT y.VendorID,y.DOLocationID,cast(y.passenger_count as
int),y.trip_distance,cast(y.total_amount as double),t.zone,t.service_zone
  FROM yellowtrip_bronze_tmp y
  INNER JOIN taxizones_view t
  ON y.DOLocationID = t.LocationID
  WHERE tpep_pickup_datetime!= tpep_dropoff_datetime and
    cast(y.passenger_count as int)!=0 and
    trip_distance!=0.0 and
    cast(y.total_amount as double)!=0.0)
```

%sql  
**describe** yellowtrip\_taxizones\_view

Table			
	col_name ▲	data_type ▲	comment ▲
1	VendorID	string	null
2	DOLocationID	string	null
3	passenger_count	int	null
4	trip_distance	float	null
5	total_amount	double	null
6	zone	string	null
7	service zone	string	null

Showing all 7 rows.

```
(spark.table("yellowtrip_taxizones_view")  
  .writeStream  
  .format("delta")  
  .option("checkpointLocation",  
f'dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/checkpoint/yellowtrip_enriched_silver")  
  .outputMode("append")  
  .table("yellowtrip_enriched_silver"))
```

```
(spark.table("yellowtrip_taxizones_view")  
  .writeStream  
  .format("delta")  
  .option("checkpointLocation", f'dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/checkpoint/yellowtrip_enriched_silver")  
  .outputMode("append")  
  .table("yellowtrip_enriched_silver"))
```

▶  36203ef2-5e52-49d6-8a29-dc7f3e429cb8 Last updated: 32 days ago

Out[22]: <pyspark.sql.streaming.query.StreamingQuery at 0x7fca8a8fda60>

**Check this stream will be continuously running**

**Check the records count of yellowtrip\_enriched\_silver delta table**

%sql  
**select** count(\*) **from** yellowtrip\_enriched\_silver

Table	
	count(1) ▲
1	97367

Showing 1 row.

## Check the file location

%fs ls dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow\_tripdata

Table				
	path	name	size	modificationTime
1	dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_05.csv	yellow_tripdata_2019_05.csv	4633703	1668673115000
2	dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_06.csv	yellow_tripdata_2019_06.csv	4628092	1668678460000

Showing all 2 rows.

Ingest another Yellow Trip data files into DBFS in the same location

dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow\_tripdata

## Check the file location

%fs ls dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow\_tripdata

Table				
	path	name	size	modificationTime
1	dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_01.csv	yellow_tripdata_2019_01.csv	1341900	1668680566000
2	dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_05.csv	yellow_tripdata_2019_05.csv	4633703	1668673115000
3	dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_06.csv	yellow_tripdata_2019_06.csv	4628092	1668678460000

Showing all 3 rows.

Re-check the records count of yellowtrip\_enriched\_silver delta table

%sql

SELECT COUNT(\*) FROM yellowtrip\_enriched\_silver

Table	
	count(1)
1	112064

Showing 1 row.



## Gold Table: Total Count of Passengers for each Zone, each vendorid

Here we read a stream of data from `yellowtrip_enriched_silver` and write another stream to create an aggregate gold table of **Total Count of Passengers for each Zone, each vendorid**.

```
(spark.readStream
  .table("yellowtrip_enriched_silver")
  .createOrReplaceTempView("yellowtrip_enriched_silver_temp"))
```

```
%sql
CREATE OR REPLACE TEMP VIEW zone_passenger_count_view AS (
  SELECT Zone, VendorID, sum(passenger_count) as total_passenger_count
  FROM yellowtrip_enriched_silver_temp
  GROUP BY Zone, VendorID)
```

```
%sql
describe zone_passenger_count_view
```

Table				
	col_name	data_type	comment	
1	Zone	string	null	
2	VendorID	string	null	
3	total_passenger_count	bigint	null	
Showing all 3 rows.				

Note that we're using `.trigger(availableNow=True)` below. This provides us the ability to continue to use the strengths of Structured Streaming while triggering this job one-time to process all available data in micro-batches.

To recap, these strengths include:

- exactly once end-to-end fault tolerant processing
- automatic detection of changes in upstream data sources

If we know the approximate rate at which our data grows, we can appropriately size the cluster we schedule for this job to ensure fast, cost-effective processing.

The customer will be able to evaluate how much updating this final aggregate view of their data costs and make informed decisions about how frequently this operation needs to be run.

Downstream processes subscribing to this table do not need to re-run any expensive aggregations.

Rather, files just need to be de-serialized and then queries based on included fields can quickly be pushed down against this already-aggregated source.

```
(spark.table("zone_passenger_count_view")
  .writeStream
  .format("delta")
  .outputMode("complete")
  .option("checkpointLocation",
    f"dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/checkpoint/zone_passenger_count_gold")
  .table("zone_passenger_count_gold"))
```

```
(spark.table("zone_passenger_count_view")
  .writeStream
  .format("delta")
  .outputMode("complete")
  .option("checkpointLocation", f"dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/checkpoint/zone_passenger_count_gold")
  .table("zone_passenger_count_gold"))
```

▶  ae1cfbd7-5013-4620-a59e-ced8a26cfe02    Last updated: 32 days ago

Out[31]: <pyspark.sql.streaming.query.StreamingQuery at 0x7fca8b9fb6a0>

Check this cell command will be continuously running

## Important Considerations for complete Output with Delta

When using `complete` output mode, we rewrite the entire state of our table each time our logic runs. While this is ideal for calculating aggregates, we **cannot** read a stream from this directory, as Structured Streaming assumes data is only being appended in the upstream logic.

The gold Delta table we have just registered will perform a static read of the current state of the data each time we run the following query.

```
%sql
select * from zone_passenger_count_gold
```

	Zone	VendorID	total_passenger_count
1	Newark Airport	2	248
2	Chinatown	4	1
3	Van Cortlandt Village	1	24
4	Bay Terrace/Fort Totten	2	23
5	West Chelsea/Hudson Yards	2	2346
6	Kingsbridge Heights	2	45
7	Sunnyside	1	191

Showing all 578 rows.

```
%sql
select count(*) from zone_passenger_count_gold
```

	count(1)
1	578

Showing 1 row.

## Check the File Location

```
%fs ls dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata
```

	path	name	size	modificationTime
1	dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_01.csv	yellow_tripdata_2019_01.csv	1341900	1668680566000
2	dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_05.csv	yellow_tripdata_2019_05.csv	4633703	1668673115000
3	dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_06.csv	yellow_tripdata_2019_06.csv	4628092	1668678460000

Showing all 3 rows.

## Load another YellowTrip data file into DBFS

```
dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata
```

## Re-Check the File Location

```
%fs ls dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata
```

	path	name	size	modificationTime
1	dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_01.csv	yellow_tripdata_2019_01.csv	1341900	1668680566000
2	dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_02.csv	yellow_tripdata_2019_02.csv	1356159	1668681749000
3	dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_05.csv	yellow_tripdata_2019_05.csv	4633703	1668673115000
4	dbfs:/FileStore/tables/nyctaxidata/landing/yellowdata/yellow_tripdata/yellow_tripdata_2019_06.csv	yellow_tripdata_2019_06.csv	4628092	1668678460000

Showing all 4 rows.

## Check the Delta Table

%sql

**select \* from** zone\_passenger\_count\_gold

Table			
	Zone ▲	VendorID ▲	total_passenger_count ▲
1	Chinatown	4	1
2	Van Cortlandt Village	1	25
3	Sunset Park West	4	2
4	Kingsbridge Heights	2	53
5	Sunnyside	1	234
6	Newark Airport	2	253
7	Bav Terrace/Fort Totten	2	23

Showing all 589 rows.

## Check the count from Delta Table after data ingestion

%sql

**select count(\*) from** zone\_passenger\_count\_gold

Table	
	count(1) ▲
1	589

Showing 1 row.

## Do it Yourself :-

Generate multi-hop data pipeline using Green Trip datasets and TaxiZones dataset to generate

1. Total Passenger count for each zone and each vendor
2. Total Trip distance for each service\_zone and each vendor
3. Total Fare amount for each borough and each vendor

## Challenge Yourself

### Implementing Delta Live Table

## Lab: Migrating SQL Notebooks to Delta Live Tables

### Learning Objectives

By the end of this lab, you should be able to:

- **Convert existing data pipelines to Delta Live Tables**

### Datasets Used

This demo uses New York taxi data sets.

The schema of our two datasets is represented below.

Note that we will be manipulating these schemas during various steps.

### Yellow Trip Data Sets Metadata with Data Type ( 18 columns)

The main dataset uses records from Yellow trip datasets in the CSV format.

<i>Column Name</i>	<i>Data type</i>
VendorID	integer
tpep_pickup_datetime	timestamp
tpep_dropoff_datetime	timestamp
passenger_count	integer
trip_distance	double
RatecodeID	integer
store_and_fwd_flag	string
PULocationID	integer
DOLocationID	integer
payment_type	integer
fare_amount	double
extra	double
mta_tax	double
tip_amount	double
tolls_amount	double
improvement_surcharge	double
total_amount	double
congestion_surcharge	double

These data will later be joined with a static table of Taxi Zones information stored in an external system to identify **Zone**, **Borough** and **Service\_Zone** by name.

### Taxi Zones Data Set Metadata with Data Type ( 4 columns)

Column Name	Data type
LocationID	integer
Borough	string
Zone	string
service_zone	string

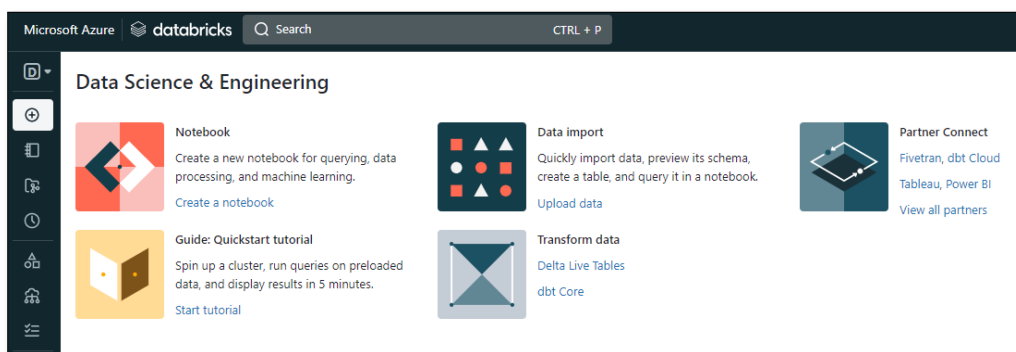
## Getting Started

### Land Initial Data

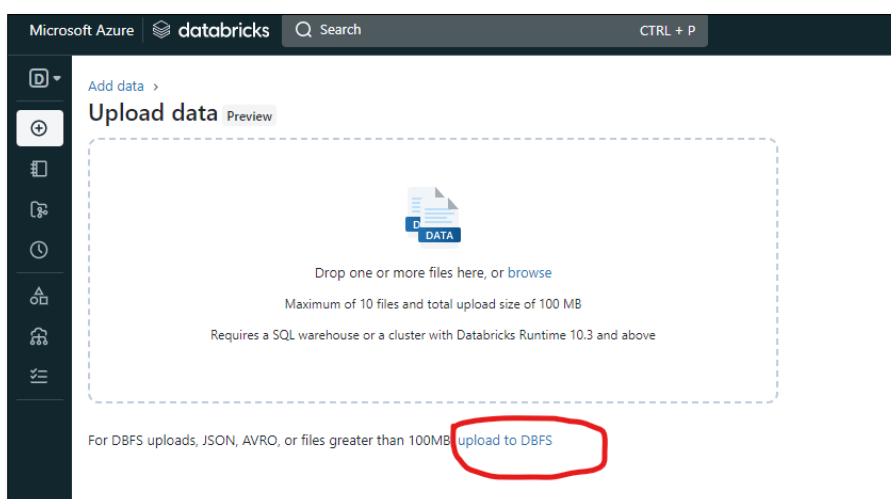
Seed the landing zone with more data before proceeding.

You will re-run this command to land additional data later.

Click on "Upload data" under "Data import" menu -



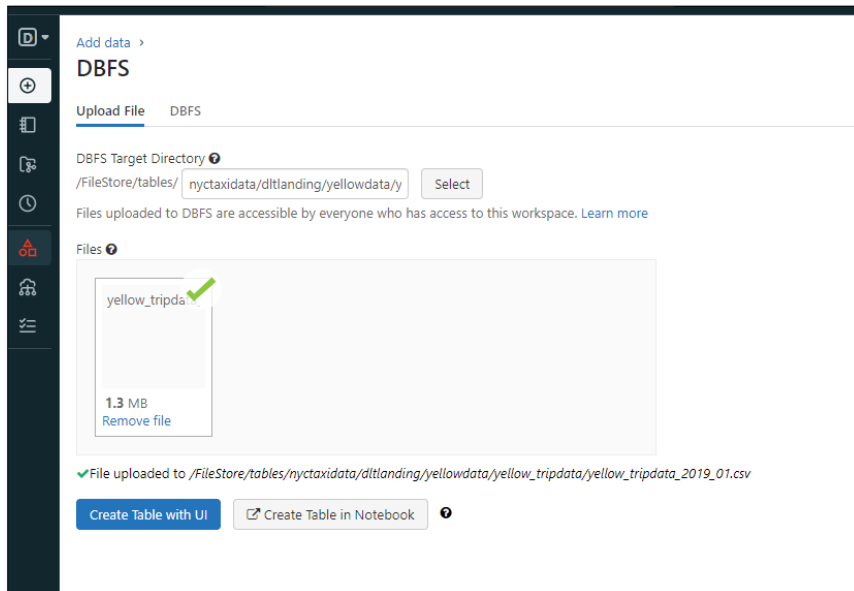
### Upload data file



## Upload Yellow Trip datasets

Landing location of Source Data in DBFS for Delta Live Table

**dbfs:/FileStore/tables/nyctaxidata/dltlanding/yellowdata/yellow\_tripdata**

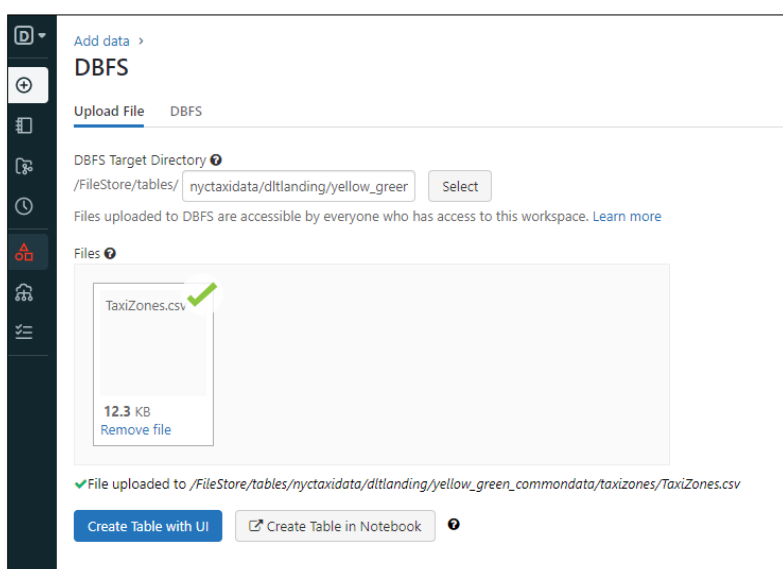


## Upload Taxi Zones datasets

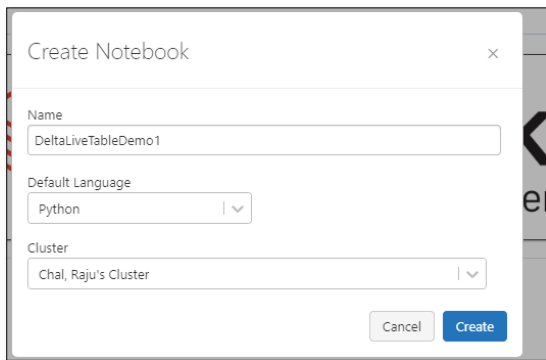
Refresh the page

Landing location of Source Data in DBFS for Delta Live Table

**dbfs:/FileStore/tables/nyctaxidata/dltlanding/yellow\_green\_commonddata/taxizones/**



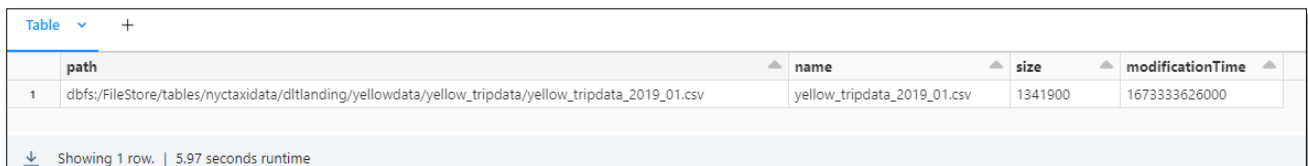
## Create a New Note Book



A dialog box titled "Create Notebook" with a close button (X) in the top right corner. It contains three input fields: "Name" with the text "DeltaLiveTableDemo1", "Default Language" with a dropdown menu showing "Python", and "Cluster" with a dropdown menu showing "Chal, Raju's Cluster". At the bottom right are "Cancel" and "Create" buttons.

Execute the following cell to check the files that has been uploaded into DBFS.

```
%fs ls dbfs:/FileStore/tables/nyctaxidata/dltlanding/yellowdata/yellow_tripdata
```



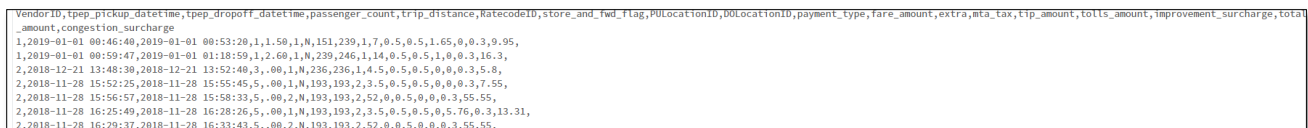
A table view showing the contents of a directory. The table has columns: path, name, size, and modificationTime. One row is visible, representing a CSV file.

	path	name	size	modificationTime
1	dbfs:/FileStore/tables/nyctaxidata/dltlanding/yellowdata/yellow_tripdata/yellow_tripdata_2019_01.csv	yellow_tripdata_2019_01.csv	1341900	1673333626000

Showing 1 row. | 5.97 seconds runtime

```
%fs head
```

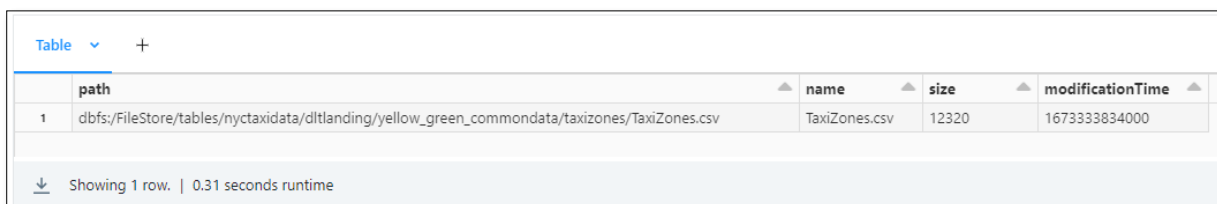
```
dbfs:/FileStore/tables/nyctaxidata/dltlanding/yellowdata/yellow_tripdata/yellow_tripdata_2019_01.csv
```



The first few lines of the CSV file, showing columns: VendorID, tpep\_pickup\_datetime, tpep\_dropoff\_datetime, passenger\_count, trip\_distance, RatecodeID, store\_and\_fwd\_flag, PULocationID, DOLocationID, payment\_type, fare\_amount, extra, mta\_tax, tip\_amount, tolls\_amount, improvement\_surcharge, total\_amount, congestion\_surcharge.

```
VendorID,tpep_pickup_datetime,tpep_dropoff_datetime,passenger_count,trip_distance,RatecodeID,store_and_fwd_flag,PULocationID,DOLocationID,payment_type,fare_amount,extra,mta_tax,tip_amount,tolls_amount,improvement_surcharge,total
amount,congestion_surcharge
1,2018-01-01 00:45:40,2018-01-01 00:53:20,1,1.50,1,N,153,239,1,7,0.5,0.5,1.65,0.0,3,9.95,
1,2019-01-01 00:50:47,2019-01-01 01:18:59,1,2.60,1,N,239,246,1,14,0.5,0.5,1,0.0,3,16.3,
2,2018-12-21 13:48:30,2018-12-21 13:52:40,3,.00,1,N,236,236,1,4,5,0.5,0.5,0.0,0.3,7.55,
2,2018-11-28 15:52:25,2018-11-28 15:55:45,5,.00,1,N,193,193,2,3,5,0.5,0.5,0.0,0.3,7.55,
2,2018-11-28 15:56:57,2018-11-28 15:58:33,5,.00,2,N,193,193,2,52,0.5,0.5,0.0,0.3,55.55,
2,2018-11-28 16:25:49,2018-11-28 16:28:26,5,.00,1,N,193,193,2,3,5,0.5,0.5,0.5,76,0.3,13.31,
2,2018-11-28 16:29:37,2018-11-28 16:33:43,5,.00,2,N,193,193,2,52,0.5,0.5,0.0,0.3,55.55,
```

```
%fs ls dbfs:/FileStore/tables/nyctaxidata/dltlanding/yellow_green_commdata/taxizones/
```



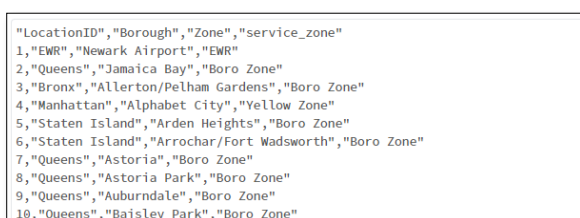
A table view showing the contents of a directory. The table has columns: path, name, size, and modificationTime. One row is visible, representing a CSV file.

	path	name	size	modificationTime
1	dbfs:/FileStore/tables/nyctaxidata/dltlanding/yellow_green_commdata/taxizones/TaxiZones.csv	TaxiZones.csv	12320	1673333834000

Showing 1 row. | 0.31 seconds runtime

```
%fs head
```

```
dbfs:/FileStore/tables/nyctaxidata/dltlanding/yellow_green_commdata/taxizones/TaxiZones.csv
```

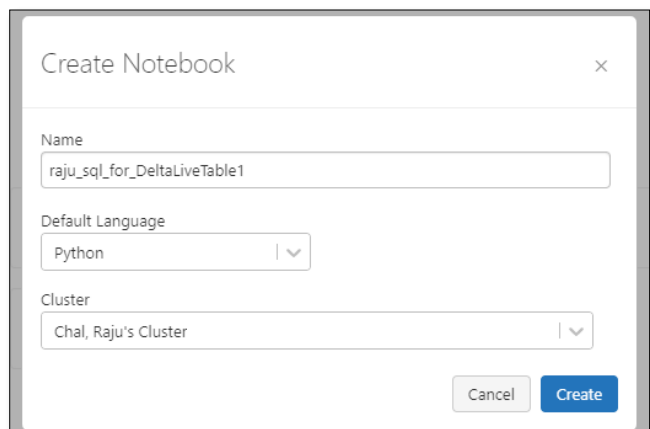


The first few lines of the CSV file, showing columns: LocationID, Borough, Zone, service\_zone.

```
"LocationID","Borough","Zone","service_zone"
1,"EWR","Newark Airport","EWR"
2,"Queens","Jamaica Bay","Boro Zone"
3,"Bronx","Allerton/Pelham Gardens","Boro Zone"
4,"Manhattan","Alphabet City","Yellow Zone"
5,"Staten Island","Arden Heights","Boro Zone"
6,"Staten Island","Arrochar/Fort Wadsworth","Boro Zone"
7,"Queens","Astoria","Boro Zone"
8,"Queens","Astoria Park","Boro Zone"
9,"Queens","Auburndale","Boro Zone"
10,"Queens","Baisley Park","Boro Zone"
```

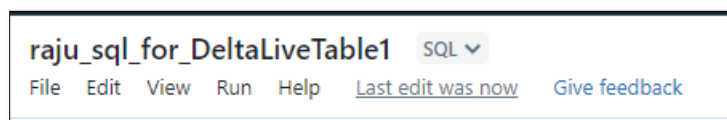


## Create another new notebook



A dialog box titled "Create Notebook" with a close button (X) in the top right corner. It contains three input fields: "Name" with the text "raju\_sql\_for\_DeltaLiveTable1", "Default Language" with a dropdown menu showing "Python", and "Cluster" with a dropdown menu showing "Chal, Raju's Cluster". At the bottom right, there are two buttons: "Cancel" and "Create".

## Make Sure default language in SQL



A header bar for a notebook. It displays the notebook name "raju\_sql\_for\_DeltaLiveTable1" followed by a language dropdown menu currently set to "SQL". Below the name, there is a menu bar with "File", "Edit", "View", "Run", and "Help". To the right of the menu bar are two links: "Last edit was now" and "Give feedback".

## Lab: Migrating a SQL Pipeline to Delta Live Tables

This notebook will be completed by you to implement a DLT pipeline using SQL.

It is **not intended** to be executed interactively, but rather to be deployed as a pipeline once you have completed your changes.

To aid in completion of this Notebook, please refer to the [DLT syntax documentation](#).

### File Locations:-

### Declare Bronze Table

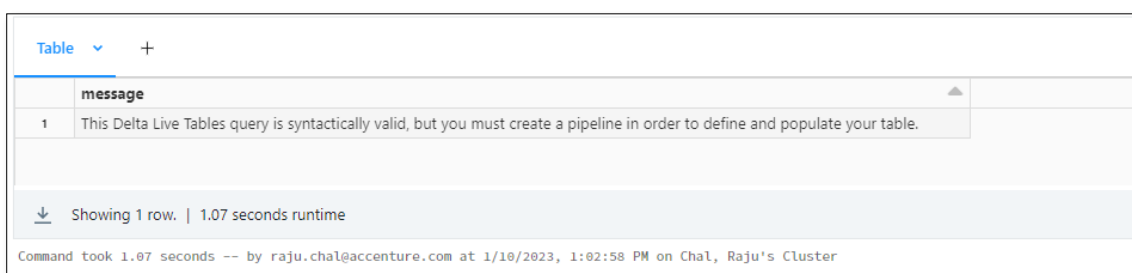
Declare a bronze table, **yellowTrip\_bronze** that ingests CSV data incrementally (using Auto Loader) from the simulated cloud source.

The source location is already supplied as an argument; using this value is illustrated in the cell below.

As we did previously, include two additional columns:

- **receipt\_time** that records a timestamp as returned by **current\_timestamp()**
- **source\_file** that is obtained by **input\_file\_name()**

```
CREATE OR REFRESH STREAMING LIVE TABLE yellowTrip_bronze
AS SELECT current_timestamp() receipt_time, input_file_name() source_file, *
FROM cloud_files("dbfs:/FileStore/tables/nyctaxidata/dltlanding/yellowdata/yellow_tripdata", "csv",
map("cloudFiles.schemaHints", "passenger_count INTEGER,trip_distance DOUBLE, total_amount
DOUBLE","header", "true", "cloudFiles.inferColumnTypes", "true"))
```



	message
1	This Delta Live Tables query is syntactically valid, but you must create a pipeline in order to define and populate your table.

Showing 1 row. | 1.07 seconds runtime

Command took 1.07 seconds -- by raju.chal@accenture.com at 1/10/2023, 1:02:58 PM on Chal, Raju's Cluster

### Taxi Zones File

Using a similar CTAS syntax, create a live **table** into the CSV data found in the *Taxi Zones* dataset.

To properly configure Auto Loader for this source, you will need to specify the following additional parameters:

option	value
header	true
cloudFiles.inferColumnTypes	true

Auto Loader configurations for CSV can be found [here](#).

```
CREATE OR REFRESH STREAMING LIVE TABLE taxi_zones
AS SELECT * FROM cloud_files("${datasets_path}/yellow_green_commondata/taxizones", "csv",
map("header", "true", "cloudFiles.inferColumnTypes", "true"))
```

## Declare Silver Tables

Our silver table, **yelloTrip\_taxiZones\_Silver**, will consist of the following fields:

```
VendorID, tpep_pickup_datetime, tpep_dropoff_datetime, passenger_count
, trip_distance, DOLocationID, total_amount, Borough, Zone, service_zone
```

```
CREATE OR REFRESH STREAMING LIVE TABLE yelloTrip_taxiZones_Silver
(CONSTRAINT positive_passenger_count EXPECT (passenger_count > 0) ON VIOLATION DROP ROW,
CONSTRAINT positive_trip_distance EXPECT (trip_distance > 0.0) ON VIOLATION DROP ROW,
CONSTRAINT positive_total_amount EXPECT (total_amount > 0.0) ON VIOLATION DROP ROW,
CONSTRAINT pickup_drop_datetime_not_same EXPECT (pickup_datetime!=dropoff_datetime ) ON
VIOLATION DROP ROW
)
AS SELECT
a.VendorID,
a.tpep_pickup_datetime pickup_datetime,
a.tpep_dropoff_datetime dropoff_datetime,
CAST(a.passenger_count AS INTEGER) passenger_count,
CAST(a.trip_distance AS DOUBLE) trip_distance,
a.DOLocationID,
CAST(a.total_amount AS DOUBLE) total_amount,
b.Borough,
b.Zone,
b.service_zone
FROM STREAM(live.yellowTrip_bronze) a
INNER JOIN STREAM(live.taxi_zones) b
ON a.DOLocationID = b.LocationID
```

Table +	
message	
1	This Delta Live Tables query is syntactically valid, but you must create a pipeline in order to define and populate your table.
Showing 1 row.   0.22 seconds runtime	

## Gold Table

Create a gold table, `vendor_zone_passenger_count`, that aggregates `passenger_count` by **VendorID** and **Zone** and delivers the following columns:

1

```
CREATE OR REFRESH LIVE TABLE vendor_zone_passenger_count
COMMENT "Total Passenger count for each yellow taxi vendor for each zone"
AS SELECT VendorID, Zone, sum(passenger_count) total_passenger_count
FROM live.yelloTrip_taxiZones_Silver
GROUP BY VendorID, Zone
```

Table +	
	message
1	This Delta Live Tables query is syntactically valid, but you must create a pipeline in order to define and populate your table.

## Create and Configure a Pipeline

### Properties to Configure the Pipeline

Pipeline Name: `DLT-Demo-rajuchal-10-jan-2023`

Target: `raju_chal_dlt_demo_10_jan_2023`

Storage Location: `dbfs:/FileStore/tables/nyctaxidata/dltphase1/dlt_demo_10_jan_2023/storage`

Notebook Path: `/Repos/raju.chal@accenture.com/data-engineering-with-databricks/08 - Delta Live Tables/DE 8.1 - DLT/DE 8.1.2 - SQL for Delta Live Tables`

Datasets Path: `dbfs:/FileStore/tables/nyctaxidata/dltlanding`

## Create and Configure a Pipeline using Web Interface

In this section you will create a pipeline using a notebook provided with the courseware. We'll explore the contents of the notebook in the following lesson.

1. Click the **Workflows** button on the sidebar.
2. Select the **Delta Live Tables** tab.
3. Click **Create Pipeline**.
4. Leave **Product Edition** as **Advanced**.
5. Fill in a **Pipeline Name** - because these names must be unique, we suggest using the **Pipeline Name** provided in the cell above.

6. For **Notebook Libraries**, use the navigator to locate and select the notebook specified above.
  - Even though this document is a standard Databricks Notebook, the SQL syntax is specialized to DLT table declarations.
  - We will be exploring the syntax in the exercise that follows.
7. Towards the bottom of the page, there is a drop down titled **Advanced**. Click on that, then:
  - Click **Add configuration**, set the "key" to **spark.master** and the "value" to **local[\*]**.
  - Click **Add configuration**, set the "key" to **datasets\_path** and the "value" to the value provided in the cell above.
8. In the **Target** field, enter the database name provided in the cell above. This should follow the pattern `<name>_<hash>_dbacademy_dewd_dlt_demo_81`
  - This field is optional; if not specified, then tables will not be registered to a metastore, but will still be available in the DBFS. Refer to the [documentation](#) for more information on this option.
9. In the **Storage location** field, enter the path provided in the cell above.
  - This optional field allows the user to specify a location to store logs, tables, and other information related to pipeline execution.
  - If not specified, DLT will automatically generate a directory.
10. For **Pipeline Mode**, select **Triggered**.
  - This field specifies how the pipeline will be run.
  - **Triggered** pipelines run once and then shut down until the next manual or scheduled update.
  - **Continuous** pipelines run continuously, ingesting new data as it arrives. Choose the mode based on latency and cost requirements.
11. Uncheck the **Enable autoscaling** box.
12. Set the number of **workers** to **0** (zero).
  - Along with the **spark.master** config above, this will create a **Single Node** clusters.
13. Check the **Use Photon Acceleration** box.
14. For **Channel**, select **Current**
15. For **Policy**, select the value provided in the cell above.

The fields **Enable autoscaling**, **Min Workers** and **Max Workers** control the worker configuration for the underlying cluster processing the pipeline.

Notice the DBU estimate provided, similar to that provided when configuring interactive clusters.

Finally, click **Create**.

Workflows > Delta Live Tables > Create pipeline

## Create pipeline

### General

**\* Pipeline name**

**\* Product edition**

[Help me choose](#)

**Pipeline mode**

☒ Triggered
 ☐ Continuous

**Cluster policy**

### Source code

**\* Notebook libraries**

[Add notebook library](#)

**Summary**

✓ Photon enabled

DBU / hour: 5

Workflows > Delta Live Tables > Create pipeline

## Create pipeline

### Destination

**Storage location**

**Target schema**

### Compute

**Cluster mode**

[Learn about Enhanced Autoscaling](#)

**\* Workers**

☒ Use Photon Acceleration

**Advanced**

Configuration

spark.master	local[*]	×
datasets_path	dbfs:/FileStore/tables/mytaxidata/dtfl	×

[Add configuration](#)

**Channel**

**Summary**

✓ Photon enabled

DBU / hour: 5

[Cancel](#) [Create](#)

Workflows > Delta Live Tables > Pipeline details

## DLT-Demo-rajuchal-10-jan-2023

Development Production Delete Permissions Settings Schedule Start

**Pipeline details**

**Name** DLT-Demo-rajuchal-10-jan-2023  
**Pipeline ID** 90965712-391f-4a44-a11a-d70d2dec8ffe  
**Paths** /Users/raju.chal@accenture.com/raju\_sql\_for\_DeltaLiveTable1  
**Run as** raju.chal@accenture.com

**Compute**

Cost	DBU / hour: 5
Cloud	Azure
Product edition	Advanced

A graph will be generated here once a pipeline update has started. Click "Start" to start an update.

All Info Warning Error

Q filter...

No events

## Open and Complete DLT Pipeline Notebook

Open the Notebook and, following the guidelines provided therein, fill in the cells where prompted to implement a multi-hop architecture similar to the one we worked with in the previous section.

## Run your Pipeline

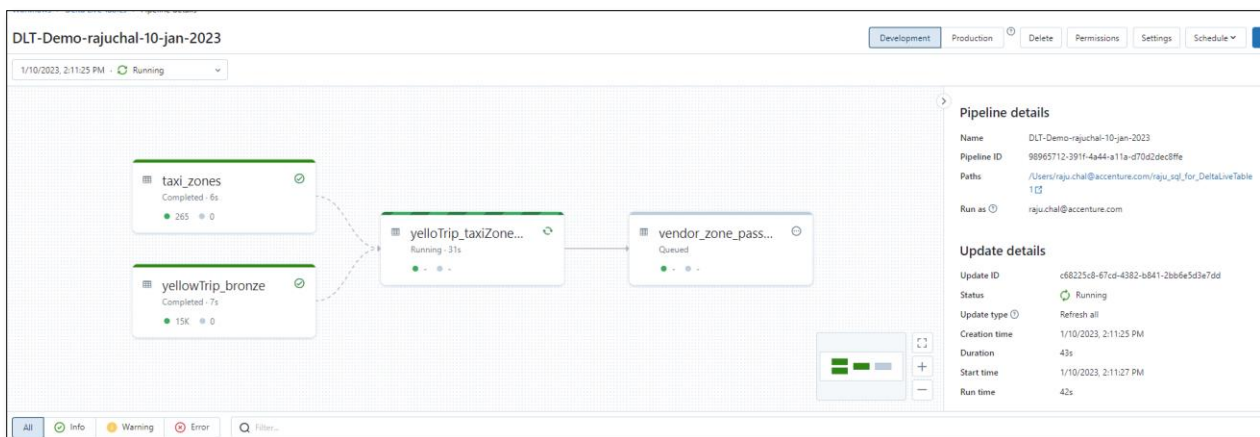
Select **Development** mode, which accelerates the development lifecycle by reusing the same cluster across runs.

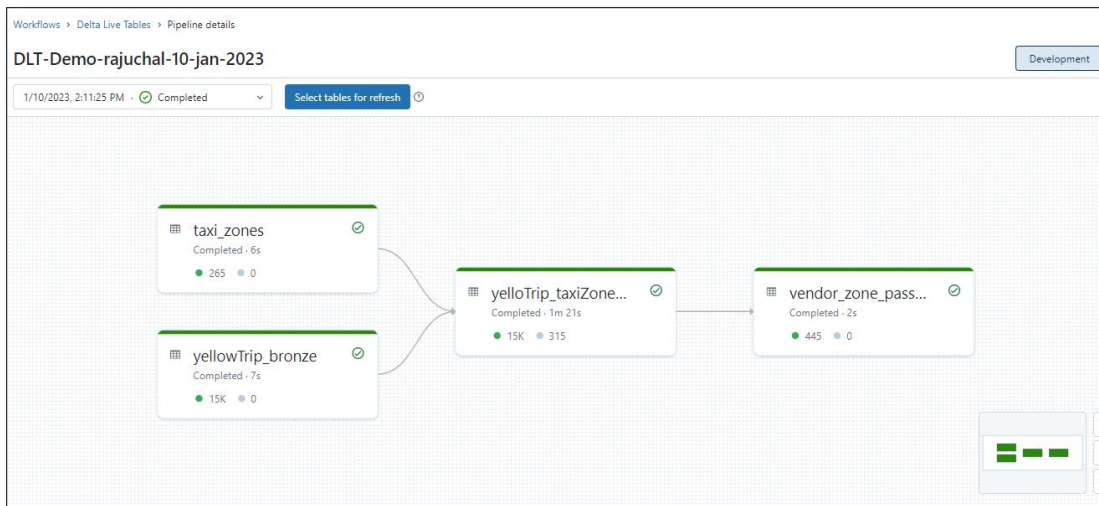
It will also turn off automatic retries when jobs fail.

Click **Start** to begin the first update to your table.

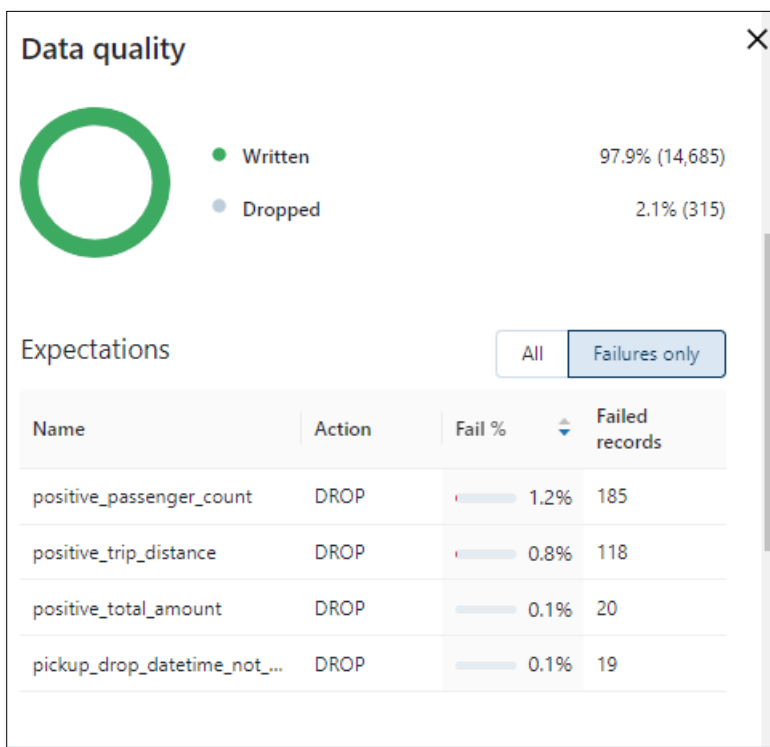

Delta Live Tables will automatically deploy all the necessary infrastructure and resolve the dependencies between all datasets.

**NOTE:** The first table update may take several minutes as relationships are resolved and infrastructure deploys.





Select yelloTrip\_taxiZones\_Silver table in data flow and check the Data Quality section

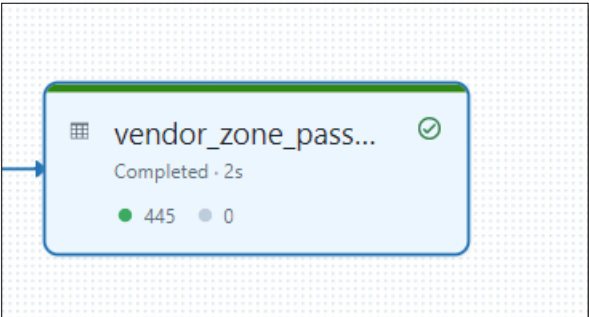




### Schema

VendorID: integer  
 pickup\_datetime: timestamp  
 dropoff\_datetime: timestamp  
 passenger\_count: integer  
 trip\_distance: double  
 DOLocationID: integer  
 total\_amount: double  
 Borough: string  
 Zone: string  
 service\_zone: string

Select vendor\_zone\_passenger\_count table in data flow and check the Data Quality section




A screenshot of a data flow task node in a workflow. The node is labeled 'vendor\_zone\_pass...' and has a green checkmark icon. Below the name, it says 'Completed - 2s'. At the bottom, there are two colored circles with numbers: a green circle with '445' and a grey circle with '0'.

### vendor\_zone\_passenger\_count

Name	vendor_zone_passenger_count
Type	Table
Path	/Users/raju.chal@accenture.com/raju_sql_for_DeltaLiveTable1
Metastore	raju_chal_dlt_demo_10_jan_2023.vendor_zone_passenger_count
Status	Completed
Start time	1/10/2023, 2:12:59 PM
Duration	2s
Comment	Total Passenger count for each yellow taxi vendor for each zone

### Data quality



A donut chart showing data quality metrics. The chart is mostly green, indicating that 100% of the data (445 rows) was written. A small grey segment indicates that 0% of the data (0 rows) was dropped.

Category	Count	Percentage
Written	445	100%
Dropped	0	0%

### Schema

VendorID: integer  
 Zone: string  
 total\_passenger\_count: long

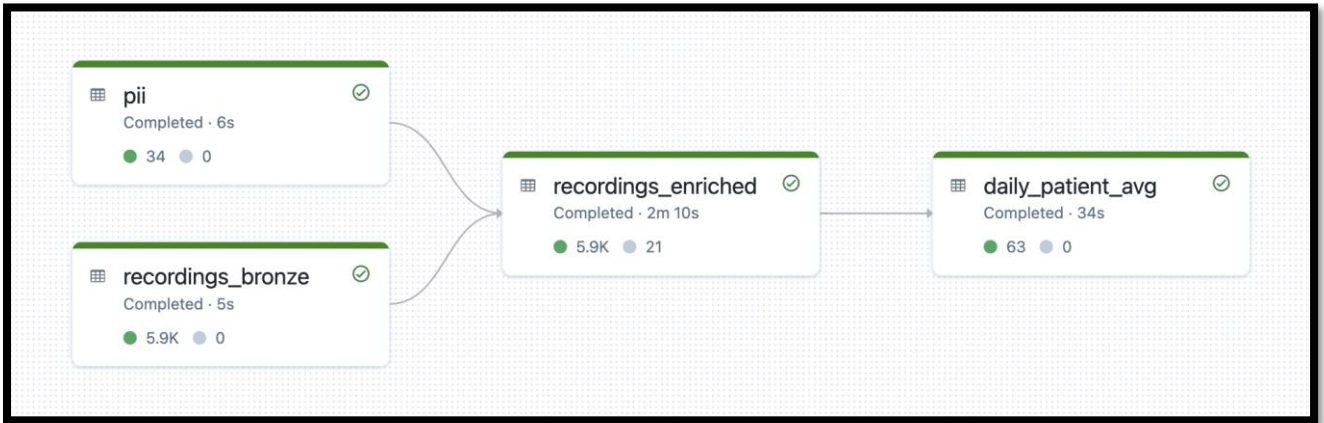
Open the DeltaLiveTableDemo1 note book that we have already created

## Troubleshooting Code in Development Mode

Don't despair if your pipeline fails the first time. Delta Live Tables is in active development, and error messages are improving all the time.

Because relationships between tables are mapped as a DAG, error messages will often indicate that a dataset isn't found.

Let's consider our DAG below:



If the error message **Dataset not found: 'recordings\_parsed'** is raised, there may be several culprits:

1. The logic defining **recordings\_parsed** is invalid
2. There is an error reading from **recordings\_bronze**
3. A typo exists in either **recordings\_parsed** or **recordings\_bronze**

The safest way to identify the culprit is to iteratively add table/view definitions back into your DAG starting from your initial ingestion tables. You can simply comment out later table/view definitions and uncomment these between runs.

## Lab: Conclusion

### Display Results

```
files = dbutils.fs.ls("dbfs:/FileStore/tables/nyctaxidata/dltphase1/dlt_demo_10_jan_2023/storage")
```

```
display(files)
```

	path	name	size	modificationTime
1	dbfs:/FileStore/tables/nyctaxidata/dltphase1/dlt_demo_10_jan_2023/storage/autoloader/	autoloader/	0	1673339758000
2	dbfs:/FileStore/tables/nyctaxidata/dltphase1/dlt_demo_10_jan_2023/storage/checkpoints/	checkpoints/	0	1673340090000
3	dbfs:/FileStore/tables/nyctaxidata/dltphase1/dlt_demo_10_jan_2023/storage/system/	system/	0	1673339737000
4	dbfs:/FileStore/tables/nyctaxidata/dltphase1/dlt_demo_10_jan_2023/storage/tables/	tables/	0	1673339994000

Showing all 4 rows. | 1.05 seconds runtime

```
files =
```

```
dbutils.fs.ls("dbfs:/FileStore/tables/nyctaxidata/dltphase1/dlt_demo_10_jan_2023/storage/system/events")
```

```
display(files)
```

	path	name	size	modificationTime
1	dbfs:/FileStore/tables/nyctaxidata/dltphase1/dlt_demo_10_jan_2023/storage/system/events/_delta_log/	_delta_log/	0	1673340837000
2	dbfs:/FileStore/tables/nyctaxidata/dltphase1/dlt_demo_10_jan_2023/storage/system/events/part-00000-0803eb17-0348-40cd-88ed-c02a5cf9e22f.c000.snappy.parquet	part-00000-0803eb17-0348-40cd-88ed-c02a5cf9e22f.c000.snappy.parquet	10729	1673340184000
3	dbfs:/FileStore/tables/nyctaxidata/dltphase1/dlt_demo_10_jan_2023/storage/system/events/part-00000-091869c2-00a0-48fd-aff7-603ebd068d5a.c000.snappy.parquet	part-00000-091869c2-00a0-48fd-aff7-603ebd068d5a.c000.snappy.parquet	10303	1673340416000
4	dbfs:/FileStore/tables/nyctaxidata/dltphase1/dlt_demo_10_jan_2023/storage/system/events/part-00000-17d680fb-97ec-4ca3-89a1-290b4a16c38a.c000.snappy.parquet	part-00000-17d680fb-97ec-4ca3-89a1-290b4a16c38a.c000.snappy.parquet	10951	1673340176000
5	dbfs:/FileStore/tables/nyctaxidata/dltphase1/dlt_demo_10_jan_2023/storage/system/events/part-00000-1d16407a-91d7-4851-b117-7674b814a06f.c000.snappy.parquet	part-00000-1d16407a-91d7-4851-b117-7674b814a06f.c000.snappy.parquet	10272	1673340296000
6	dbfs:/FileStore/tables/nyctaxidata/dltphase1/dlt_demo_10_jan_2023/storage/system/events/part-00000-33b8d0ea-46b7-450d-ad5c-c6c1131eb47a.c000.snappy.parquet	part-00000-33b8d0ea-46b7-450d-ad5c-c6c1131eb47a.c000.snappy.parquet	10242	1673339966000
↓	dbfs:/FileStore/tables/nyctaxidata/dltphase1/dlt_demo_10_jan_2023/storage/system/events/part-00000-388d7b79-1b46-4b88-ad80-34e9cdc7b27e.c000.snappy.parquet	part-00000-388d7b79-1b46-4b88-ad80-34e9cdc7b27e.c000.snappy.parquet	26639	1673339761000

Showing all 22 rows. | 0.41 seconds runtime

```
%sql
```

```
SELECT * FROM
```

```
delta.`dbfs:/FileStore/tables/nyctaxidata/dltphase1/dlt_demo_10_jan_2023/storage/system/events`
```

	id	sequence	origin	timestamp
1	59c74480-90c2-11ed-816b-5ab619acfd0a	["data_plane_id": {"instance": "execution", "seq_no": 1673339735745043}, {"control_plane_seq_no": 1673053436448576}]	["cloud": "Azure", "region": "eastus", "org_id": 6797437592858319, "user_id": null, "pipeline_id": "98965712-391f-4a44-a11a-d70d2dec8ffe", "pipeline_name": "DLT-Demo-rajuchal-10-jan-2023", "cluster_id": null, "update_id": "59955712-b747-4ef7-a15b-8ec69a47f660", "maintenance_id": null, "table_id": null, "table_name": null, "flow_id": null, "flow_name": null, "batch_id": null, "uc_resource_id": null}]	2023-01-10T
2	59c5e4f0-90c2-11ed-816b-5ab619acfd0a	["data_plane_id": {"instance": "execution", "seq_no": 1673339735745042}, {"control_plane_seq_no": 1673053436448575}]	["cloud": "Azure", "region": "eastus", "org_id": 6797437592858319, "user_id": null, "pipeline_id": "98965712-391f-4a44-a11a-d70d2dec8ffe", "pipeline_name": "DLT-Demo-rajuchal-10-jan-2023", "cluster_id": null, "update_id": "59955712-b747-4ef7-a15b-8ec69a47f660", "maintenance_id": null, "table_id": null, "table_name": null, "flow_id": null, "flow_name": null, "batch_id": null, "uc_resource_id": null}]	2023-01-10T
	613371d0-90c2-11ed-b5ee-0016e05e396	["data_plane_id": {"instance": "execution", "seq_no": 1673339735745046}, {"control_plane_seq_no": null}]	["cloud": "Azure", "region": "eastus", "org_id": 6797437592858319, "user_id": null, "pipeline_id": "98965712-391f-4a44-a11a-d70d2dec8ffe", "pipeline_name": "DLT-Demo-rajuchal-10-jan-2023", "cluster_id": "0110-083241-hg6beem", "update_id": "59955712-b747-4ef7-a15b-8ec69a47f660", "maintenance_id": null, "table_id": null, "table_name": null, "flow_id": null, "flow_name": null, "batch_id": null, "uc_resource_id": "98965712-391f-4a44-a11a-d70d2dec8ffe"]	2023-01-10T

Showing all 254 rows. | 2.08 seconds runtime

Refreshed 1 minute ago

SQL cell result stored as PySpark data frame `_sqldf`. [Learn more](#)

```
files = dbutils.fs.ls("dbfs:/FileStore/tables/nyctaxidata/dltphase1/dlt_demo_10_jan_2023/storage/tables/")

display(files)
```



	path	name	size	modificationTime
1	dbfs:/FileStore/tables/nyctaxidata/dltphase1/dlt_demo_10_jan_2023/storage/tables/taxi_zones/	taxi_zones/	0	1673340094000
2	dbfs:/FileStore/tables/nyctaxidata/dltphase1/dlt_demo_10_jan_2023/storage/tables/vendor_zone_passenger_count/	vendor_zone_passenger_count/	0	1673340180000
3	dbfs:/FileStore/tables/nyctaxidata/dltphase1/dlt_demo_10_jan_2023/storage/tables/yelloTrip_taxiZones_Silver/	yelloTrip_taxiZones_Silver/	0	1673340178000
4	dbfs:/FileStore/tables/nyctaxidata/dltphase1/dlt_demo_10_jan_2023/storage/tables/yellowTrip_bronze/	yellowTrip_bronze/	0	1673340095000

Showing all 4 rows. | 0.38 seconds runtime

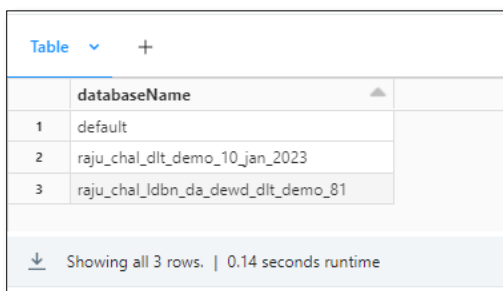
Command took 0.38 seconds -- by raju.chal@accenture.com at 1/10/2023, 2:38:15 PM on Chal, Raju's Cluster

Assuming your pipeline runs successfully, display the contents of the gold table.

**NOTE:** Because we specified a value for **Target**, tables are published to the specified database. Without a **Target** specification, we would need to query the table based on its underlying location in DBFS (relative to the **Storage Location**).

```
%sql
```

```
show databases
```

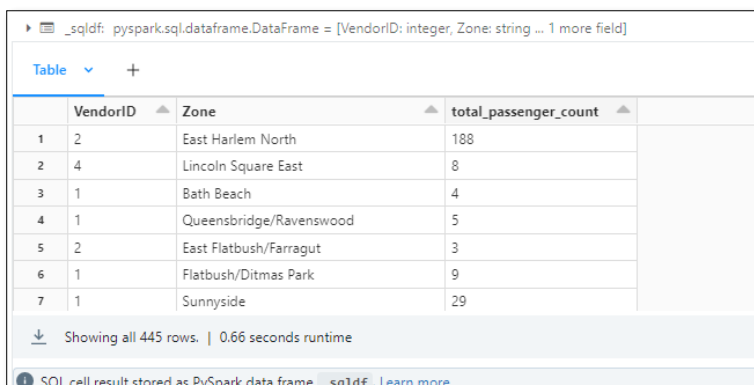


	databaseName
1	default
2	raju_chal_dlt_demo_10_jan_2023
3	raju_chal_ldbn_da_dewd_dlt_demo_81

Showing all 3 rows. | 0.14 seconds runtime

```
%sql
```

```
SELECT * FROM
raju_chal_dlt_demo_10_jan_2023.vendor_zone_passenger_count
```



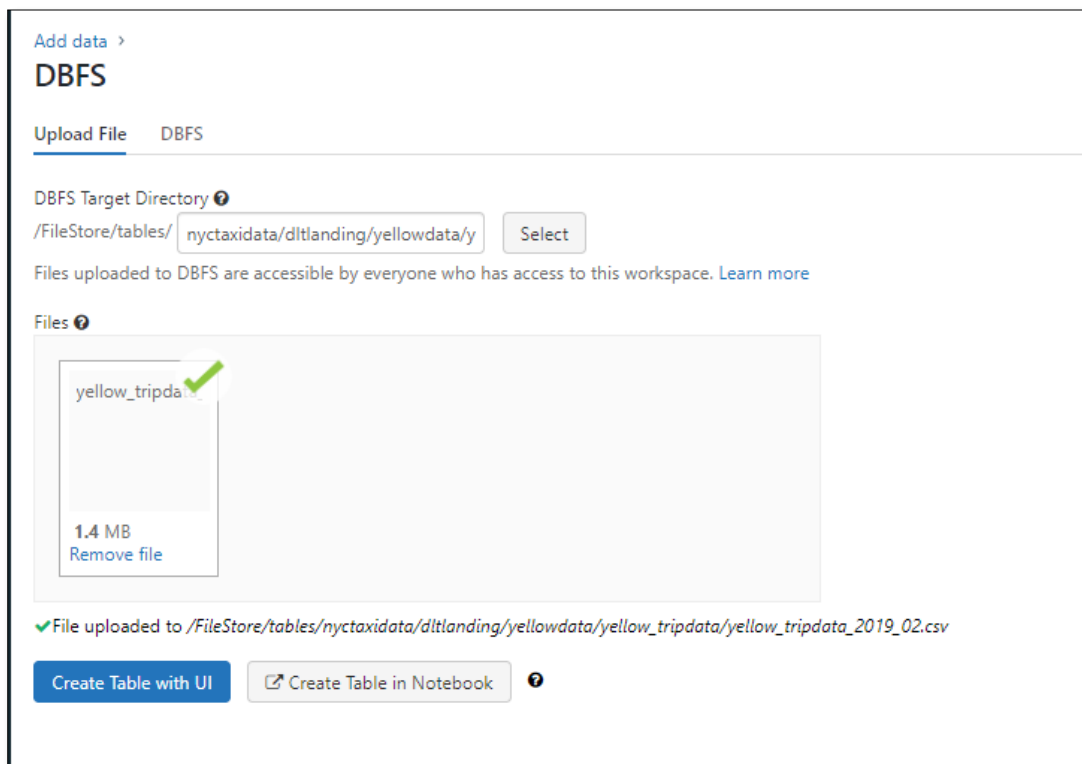
	VendorID	Zone	total_passenger_count
1	2	East Harlem North	188
2	4	Lincoln Square East	8
3	1	Bath Beach	4
4	1	Queensbridge/Ravenswood	5
5	2	East Flatbush/Farragut	3
6	1	Flatbush/Ditmas Park	9
7	1	Sunnyside	29

Showing all 445 rows. | 0.66 seconds runtime

SQL cell result stored as PySpark data frame `_sqldf`. [Learn more](#)

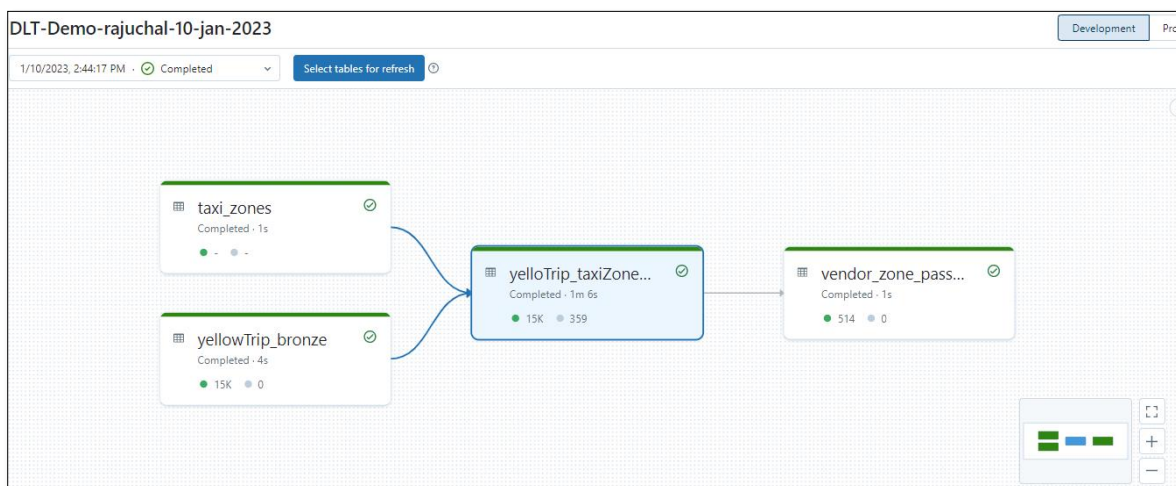
Upload another yellow trip data file into DBFS location

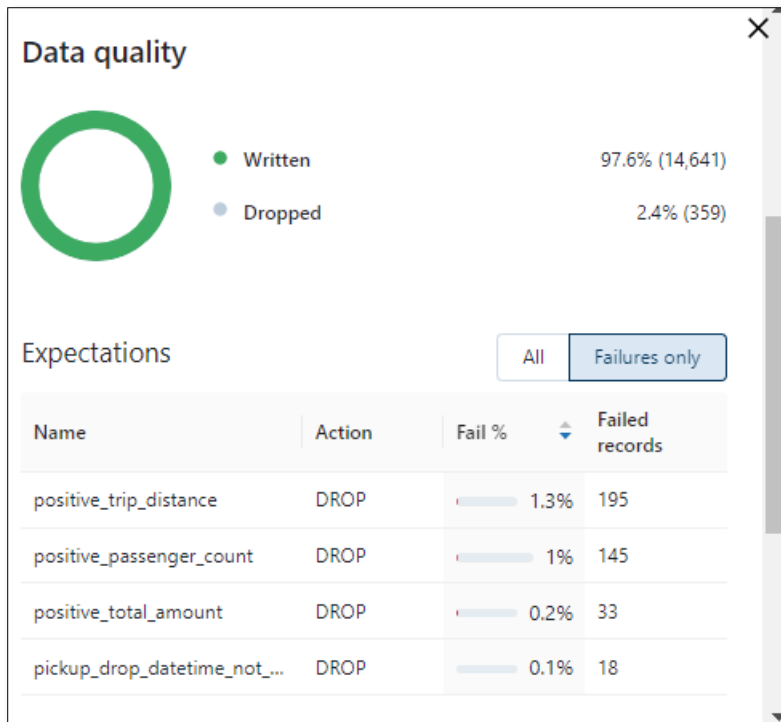
**dbfs:/FileStore/tables/nyctaxidata/dltlanding/yellowdata/yellow\_tripdata**



Feel free to run it a couple more times if desired.

Following this, run the pipeline again and view the results.





Feel free to re-run the cell above to gain an updated view of the **vendor\_zone\_passenger\_count** table.

%sql

SELECT \* FROM raju\_chal\_dlt\_demo\_10\_jan\_2023.vendor\_zone\_passenger\_count

Table +

	VendorID	Zone	total_passenger_count
1	2	East Harlem North	348
2	1	Queensbridge/Ravenswood	13
3	2	East Flatbush/Farragut	11
4	2	Bronx Park	2
5	1	Flatbush/Ditmas Park	26
6	2	Queensbridge/Ravenswood	34
7	1	Sunnyside	72

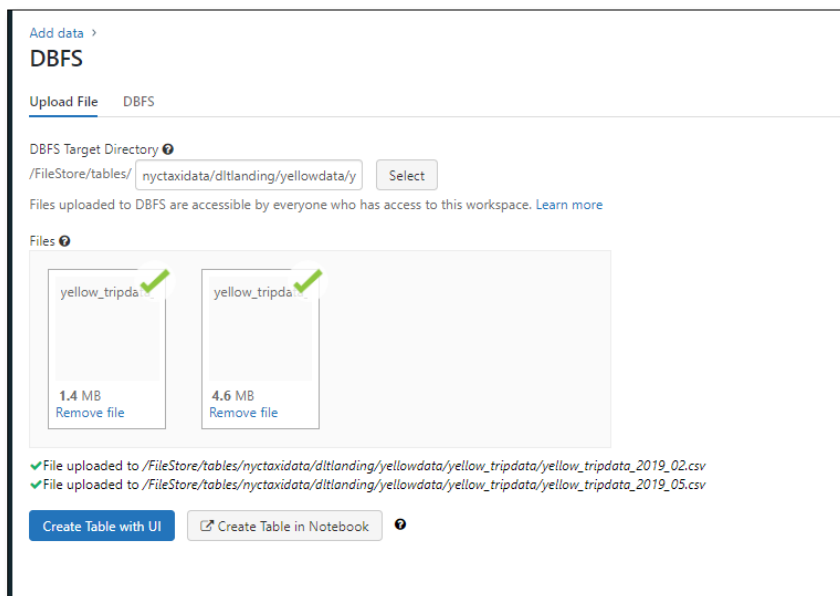
Showing all 514 rows. | 0.92 seconds runtime

SQL cell result stored as PySpark data frame `_sqldf`. [Learn more](#)

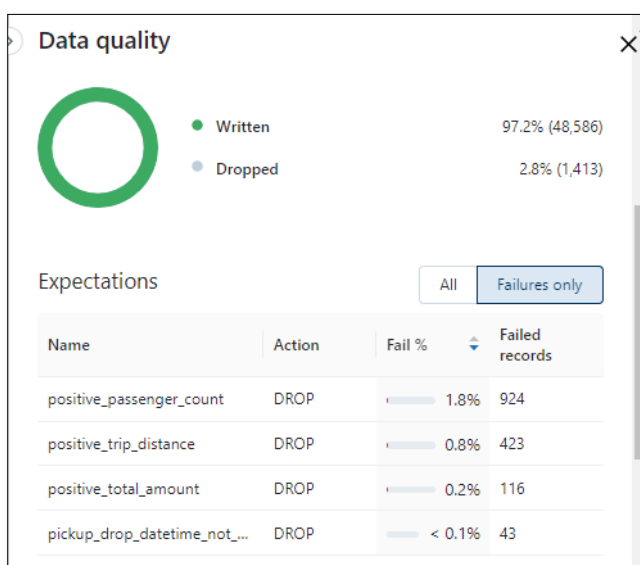
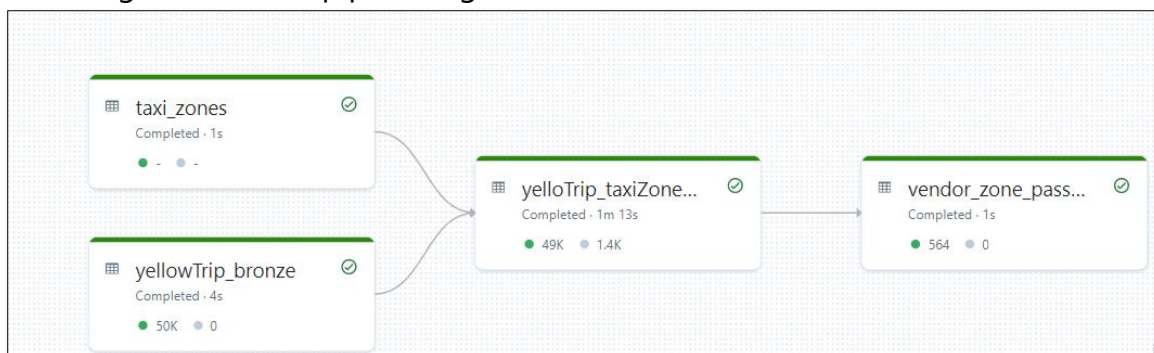
Command took 0.92 seconds -- by raju.chal@accenture.com at 1/10/2023, 2:51:06 PM on Chal, Raju's Cluster

Upload another yellow trip data file into DBFS location

dbfs:/FileStore/tables/nyctaxidata/dltlanding/yellowdata/yellow\_tripdata



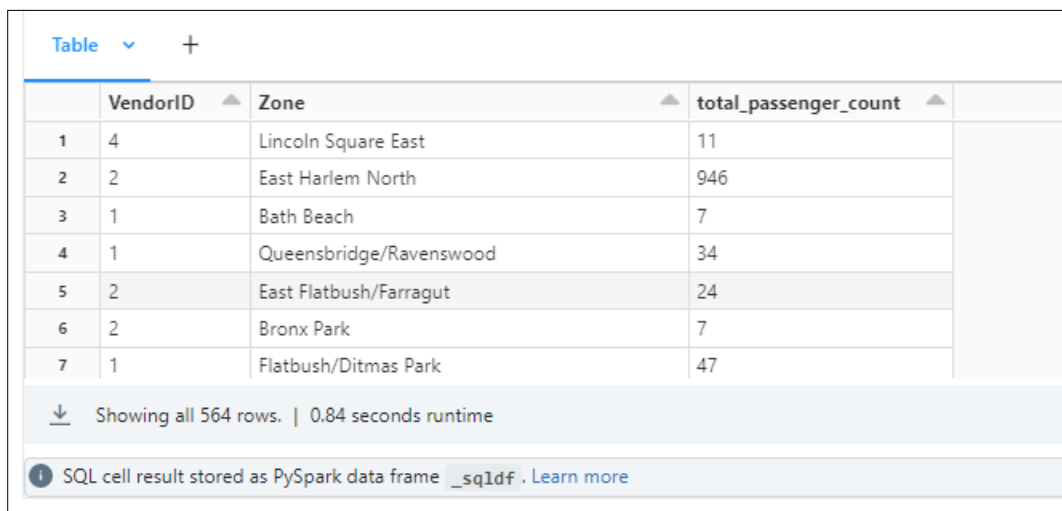
Following this, run the pipeline again and view the results.



Check the updated view of the **vendor\_zone\_passenger\_count** table.

```
%sql
```

```
SELECT * FROM raju_chal_dlt_demo_10_jan_2023.vendor_zone_passenger_count
```



	VendorID	Zone	total_passenger_count	
1	4	Lincoln Square East	11	
2	2	East Harlem North	946	
3	1	Bath Beach	7	
4	1	Queensbridge/Ravenswood	34	
5	2	East Flatbush/Farragut	24	
6	2	Bronx Park	7	
7	1	Flatbush/Ditmas Park	47	

Showing all 564 rows. | 0.84 seconds runtime

SQL cell result stored as PySpark data frame `_sqldf`. [Learn more](#)

## Summary

In this lab, you learned to convert an existing data pipeline to a Delta Live Tables SQL pipeline, and deployed that pipeline using the DLT UI.



%sql

show databases

Table ▾ +	
	databaseName ▲
1	default
2	raju_chal_dlt_demo_10_jan_2023
3	raju_chal_ldbn_da_dewd_dlt_demo_81

%sql

use raju\_chal\_dlt\_demo\_10\_jan\_2023

Table ▾ +

	database ▲	tableName ▲	isTemporary ▲
1	raju_chal_dlt_demo_10_jan_2023	taxi_zones	false
2	raju_chal_dlt_demo_10_jan_2023	vendor_zone_passenger_count	false
3	raju_chal_dlt_demo_10_jan_2023	yellotrip_taxizones_silver	false
4	raju_chal_dlt_demo_10_jan_2023	yellowtrip_bronze	false

%sql

DESCRIBE raju\_chal\_dlt\_demo\_10\_jan\_2023.yellowtrip\_bronze

Table ▾ +

	col_name ▲	data_type ▲	comment ▲	
1	receipt_time	timestamp	null	
2	source_file	string	null	
3	VendorID	int	null	
4	tpep_pickup_datetime	timestamp	null	
5	tpep_dropoff_datetime	timestamp	null	
6	passenger_count	int	null	
7	trip_distance	double	null	

[⏴](#) Showing all 21 rows. | 0.23 seconds runtime

%sql

DESCRIBE raju\_chal\_dlt\_demo\_10\_jan\_2023.yellowtrip\_taxizones\_silver

📄 \_sqldf: pyspark.sql.dataframe.DataFrame = [col\_name: string, data\_type: string ... 1 more field]

	col_name	data_type	comment
1	VendorID	int	null
2	pickup_datetime	timestamp	null
3	dropoff_datetime	timestamp	null
4	passenger_count	int	null
5	trip_distance	double	null
6	DOLocationID	int	null
7	total_amount	double	null

Showing all 10 rows. | 0.18 seconds runtime

SQL cell result stored as PySpark data frame `_sqldf`. [Learn more](#)

%sql

select \* from raju\_chal\_dlt\_demo\_10\_jan\_2023.yellowtrip\_bronze limit 10

	receipt_time	source_file	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	passenger_count	trip_distance	RatecodeID	store_
1	2023-01-10T08:41:34.005+0000	dbfs:/FileStore/tables/nyctaxidata/dt/landing/yellowdata/yellow_tripdata_2019_01.csv	1	2019-01-01T00:46:40.000+0000	2019-01-01T00:53:20.000+0000	1	1.5	1	N
2	2023-01-10T08:41:34.005+0000	dbfs:/FileStore/tables/nyctaxidata/dt/landing/yellowdata/yellow_tripdata_2019_01.csv	1	2019-01-01T00:59:47.000+0000	2019-01-01T01:18:59.000+0000	1	2.6	1	N
3	2023-01-10T08:41:34.005+0000	dbfs:/FileStore/tables/nyctaxidata/dt/landing/yellowdata/yellow_tripdata_2019_01.csv	2	2018-12-21T13:48:30.000+0000	2018-12-21T13:52:40.000+0000	3	0	1	N
4	2023-01-10T08:41:34.005+0000	dbfs:/FileStore/tables/nyctaxidata/dt/landing/yellowdata/yellow_tripdata_2019_01.csv	2	2018-11-28T15:52:25.000+0000	2018-11-28T15:55:45.000+0000	5	0	1	N
5	2023-01-10T08:41:34.005+0000	dbfs:/FileStore/tables/nyctaxidata/dt/landing/yellowdata/yellow_tripdata_2019_01.csv	2	2018-11-28T15:56:57.000+0000	2018-11-28T15:58:33.000+0000	5	0	2	N
6	2023-01-10T08:41:34.005+0000	dbfs:/FileStore/tables/nyctaxidata/dt/landing/yellowdata/yellow_tripdata_2019_01.csv	2	2018-11-28T16:25:49.000+0000	2018-11-28T16:28:26.000+0000	5	0	1	N
7	2023-01-10T08:41:34.005+0000	dbfs:/FileStore/tables/nyctaxidata/dt/landing/yellowdata/yellow_tripdata_2019_01.csv	2	2018-11-28T16:29:37.000+0000	2018-11-28T16:33:43.000+0000	5	0	2	N

Showing all 10 rows. | 1.16 seconds runtime

SQL cell result stored as PySpark data frame `_sqldf`. [Learn more](#)

Copyright © 2019 Accenture  
All rights reserved.

Accenture and its logo are  
trademarks of Accenture.

